

SoC 101:

a.k.a., "*Everything you wanted to know about a computer but were afraid to ask*"

Lecture 7: Operating Systems

Prof. Adam Teman

EnICS Labs, Bar-Ilan University

2 October 2023



Emerging Nanoscaled
Integrated Circuits and Systems Labs

The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



Lecture Overview

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

Introduction to Operating Systems

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

The Kernel

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs 12 Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

Interrupt Handling

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs 19 Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

Scheduling: Processes and Threads

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs 27 Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

Protection and Synchronization

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs 37 Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction > The Kernel > Interrupts > Scheduling > Synchronization > RTOS

RTOS

The Alexander Kofkin Faculty of Engineering Bar-Ilan University

enIcs 39 Emerging Nanoscaled Integrated Circuits and Systems Labs

```
graph TD; Applications[Applications] --> OS((OS (kernel))); Applications --> Scheduling[Scheduling]; Applications --> Protection[Protection]; Applications --> Interfacing[Interfacing]; OS --- Scheduling; OS --- Protection; OS --- Interfacing;
```

Introduction

The Kernel

Interrupts

Scheduling

Synchronization

RTOS

Introduction to Operating Systems



Emerging Nanoscaled
Integrated Circuits and Systems Labs

The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



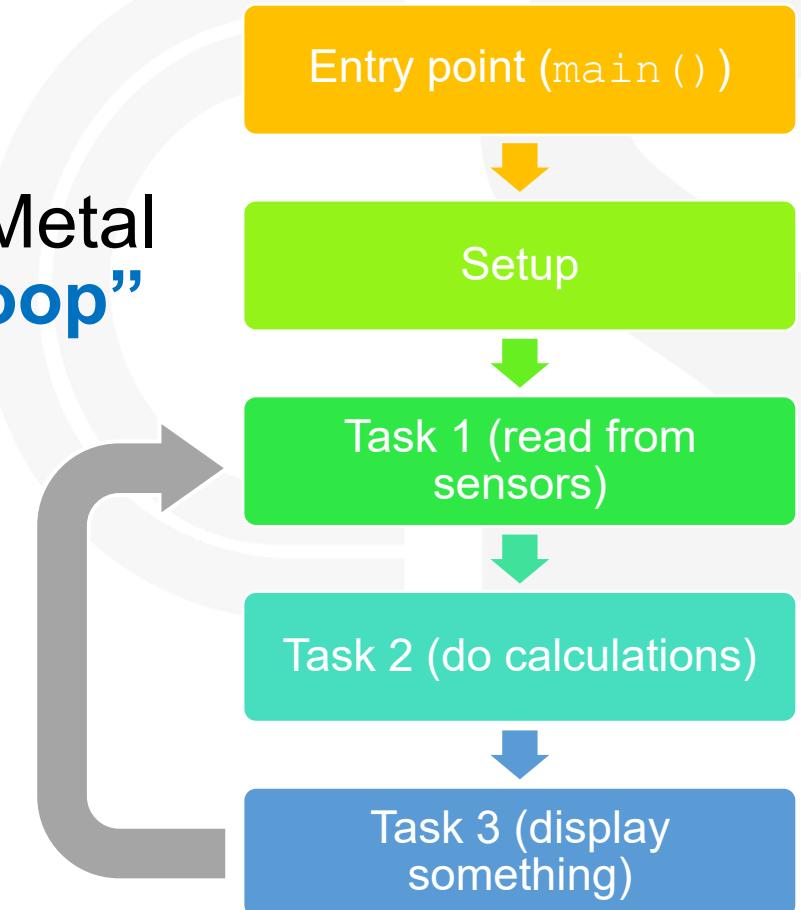
Where we started...

- What do we need in order to push a button and blink a LED?

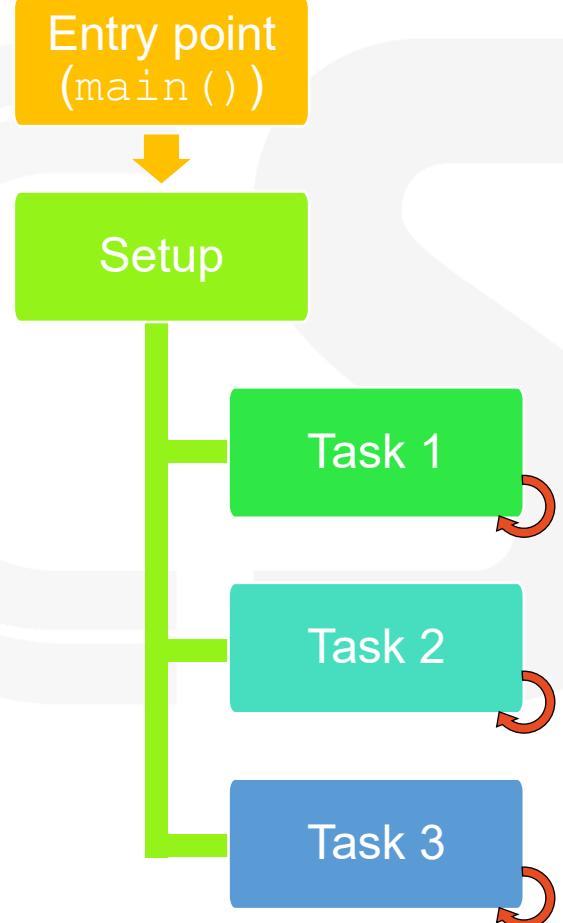
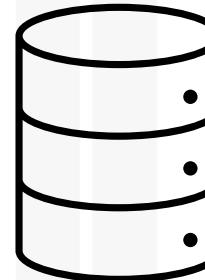
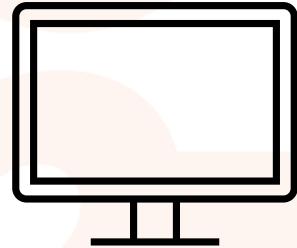
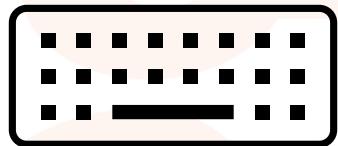


The Bare Metal “Super Loop”

- Easy to implement
- Little overhead
- Easy to debug
- Suitable for a handful of tasks

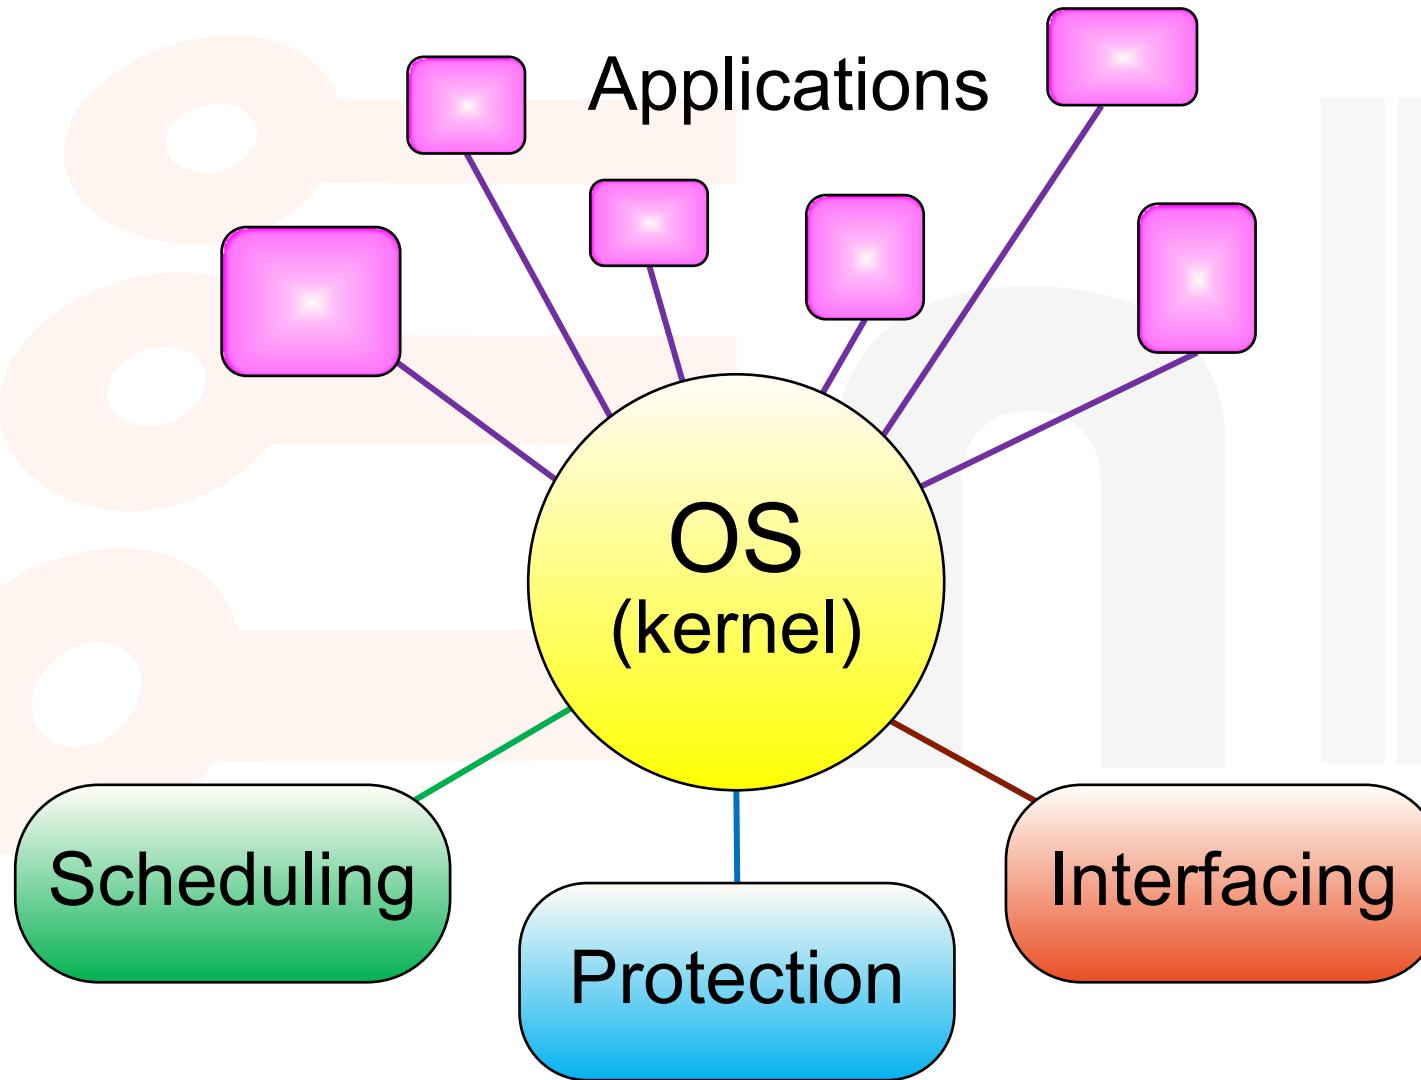


But this can quickly get complicated...



- Maybe, we need something a bit more sophisticated....

The Operating System



Referee

- Resource allocation among users, applications
- Isolation of different users, applications from each other
- Communication between users, applications



Illusionist

- Each application appears to have the entire machine to itself
- Infinite number of processors, infinite amount of memory, reliable storage, reliable network transport



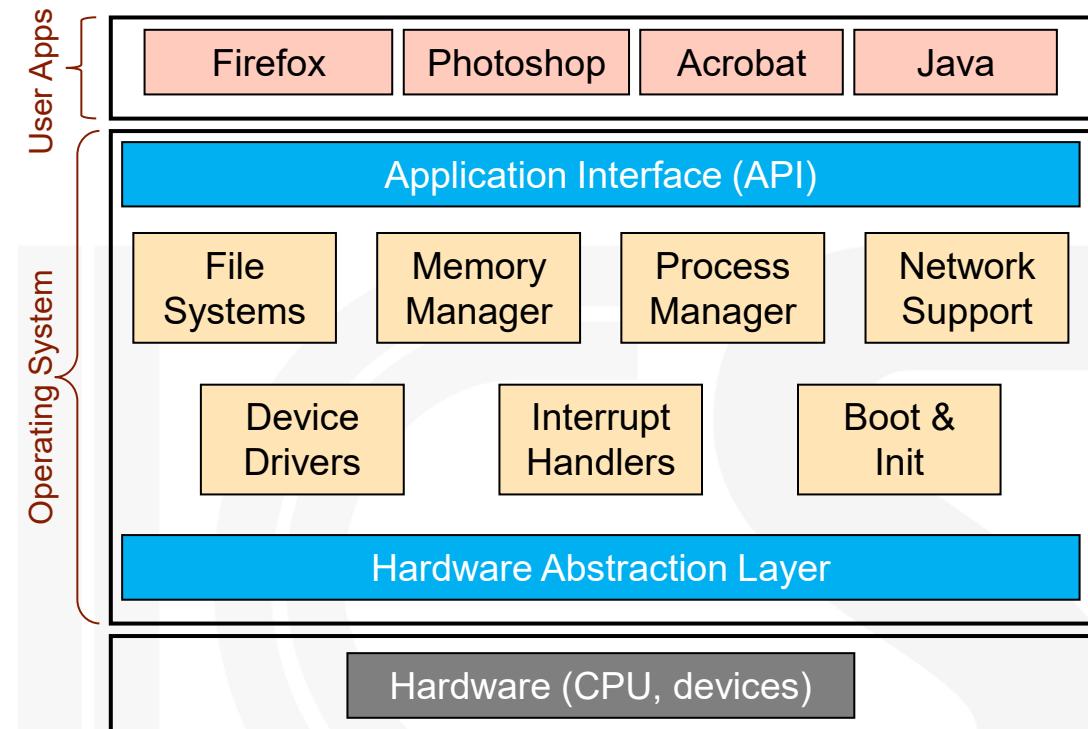
Glue

- Libraries, user interface widgets, ...
- Reduces cost of developing software



The Operating System

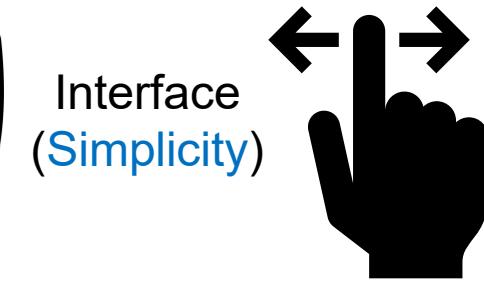
- *“The OS is everything you don’t need to write in order to run your application”*
- Primary goals of an Operating System:
 - To **simplify** program execution.
 - To use computer **hardware** efficiently.
 - To improve overall system **reliability**.
 - To provide isolation, security and **protection**.
 - To make application **software** portable and versatile.



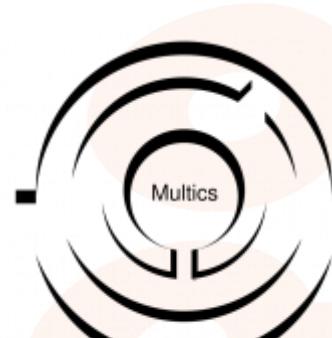
Source: Uwash, CSE 451



Protection
of resources
(Isolation)



Evolution of Operating Systems



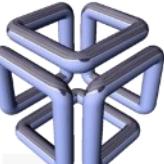
MULTICS
Bell Labs, GE, MIT
(1964-1969)



Unix
Ken Thompson,
Dennis Richie,
AT&T Bell Labs
(1969)



Unix System V
AT&T
(1974)



SGI-IRIX
Silcon Graphics



AIX
IBM



HP-UX
HP



Solaris (SunOS)
Oracle/Sun



BSD
Berkeley (1978)



NEXTSTEP



DarwinOS



Mac OS
Mac OS X (2001)
OS X (2012)
MacOS (2016)



iOS
(2007)



MS-DOS
Microsoft
(1981)



Windows 1 (1985)
Windows 95 (1995)



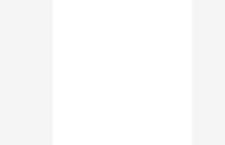
Windows XP (2001)



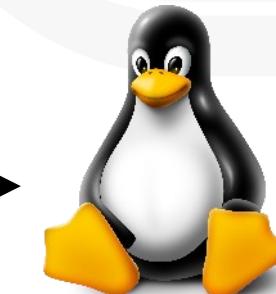
DEC (1977)



Windows NT (1993)



GNU
Richard Stallman
(1983)



Linux
Linus Torvalds
(1991)



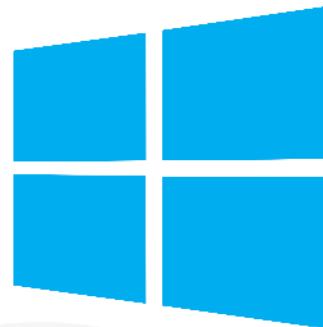
Android
(2008)

Main Operating Systems Today



ANDROID

42%



Windows®

28%



iOS

16%



Mac OS X

10%

Linux™



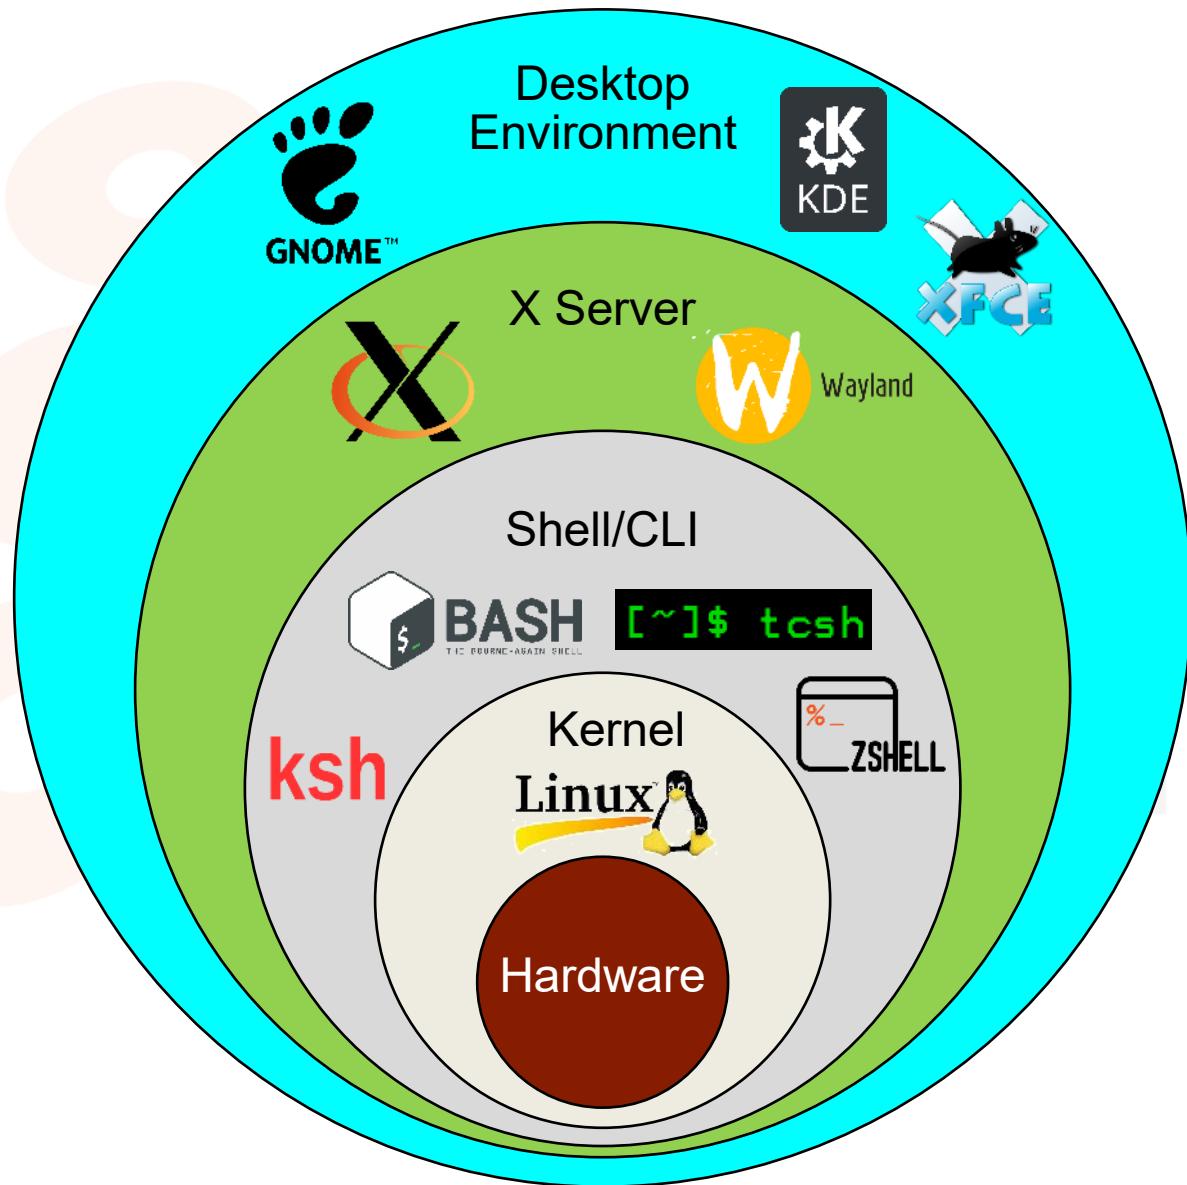
2%

Data according to Google Bard 7/2023
© Adam Teman, 2023

A Final Note: The Unix-Like Terminology

- **POSIX (Portable OS Interface)** is a standard for “**Unix-like**” Operating Systems
 - Basically, APIs that all compatible OSs understand
 - e.g., `cd`, `ls`, `fork`, `grep`, `echo`, `$HOME`, `$PATH`, etc.
- **Unix stopped being free** in 1982
 - So, Richard Stallman (MIT) started the **GNU** (“**GNU’s Not Unix!**”) project.
 - **GNU** has become a collection of free software (e.g., `GCC`, `glibc`, `GDB`, `Bash`)
- **Linux is a Kernel developed by Linus Torvalds in 1991**
 - Combining **Linux** with **GNU** provides a **fully functional and free OS**.
- **GNU/Linux are usually provided as “Distros”**
 - A **Distro** also includes the **X Server**, the **Desktop**, the **Package Manager**, etc.
 - Popular distros: **Fedora (RedHat)**, **Debian**, **Suse**, **Ubuntu**, **CentOS**, **Arch**,...

A Final Note: The Unix-Like Terminology



Distros



Package Managers

%> apt-get install	%> yum install
%> dnf install	%> pacman -S
%> zypper install	%> emerge -ask

Introduction

The Kernel

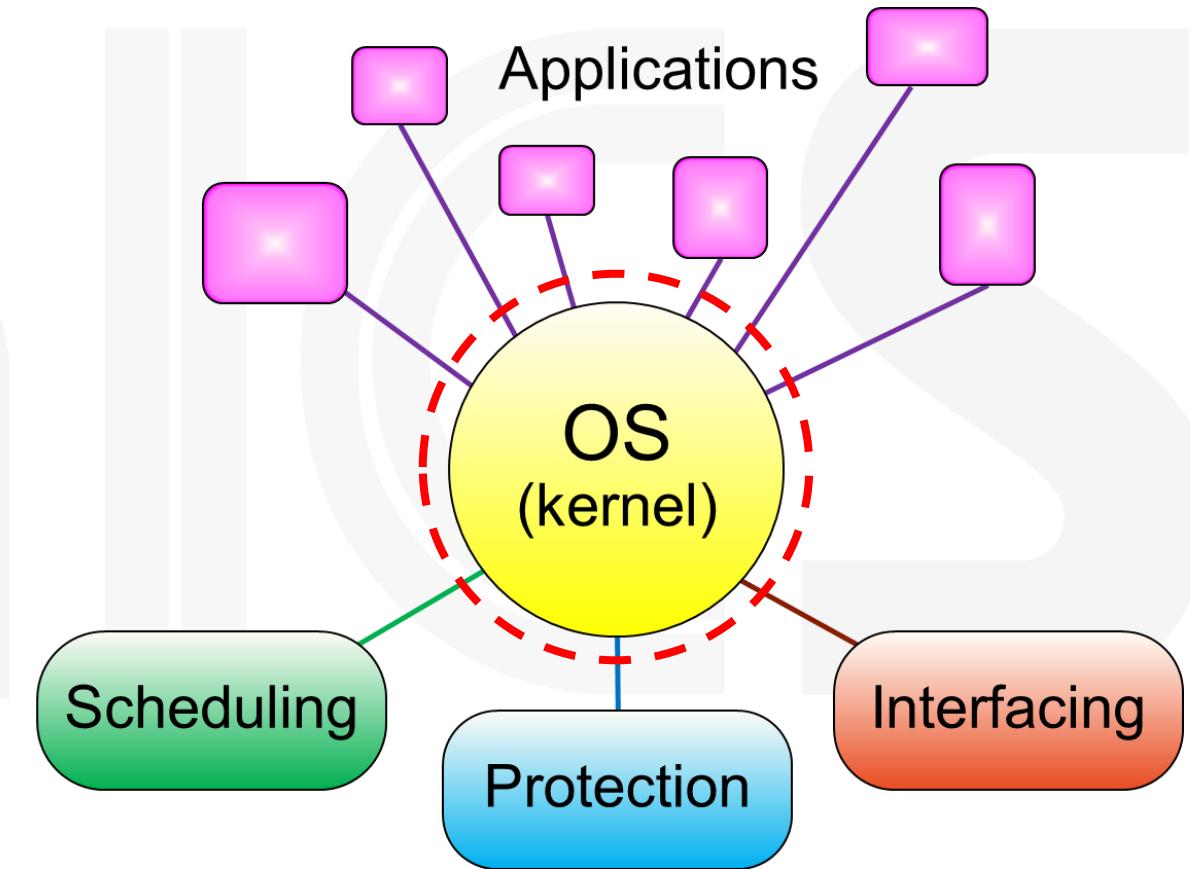
Interrupts

Scheduling

Synchronization

RTOS

The Kernel

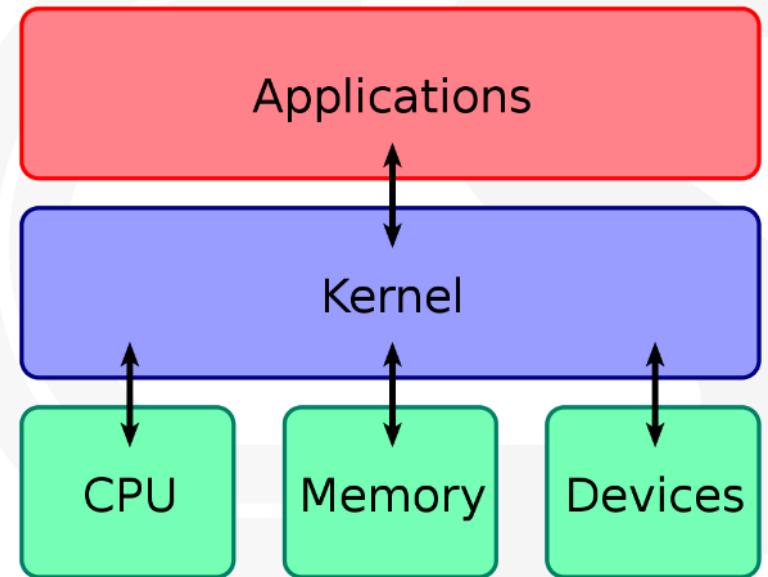


The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



The Kernel

- The **Kernel** is the **core of the Operating System**
 - The Kernel is the **first program** (after the bootloader) that is loaded into memory and run.
 - The Kernel is **always resident** in memory.
 - The Kernel is the program that **executes forever**.
- The Kernel has **complete control** over everything in the system.
 - The Kernel facilitates **interactions** between hardware and software.
 - **Everything else is an application** with respect to the Kernel.
- The Kernel is **privileged** and can carry out many operations that no other program is allowed to.



Source: Wikipedia

Hardware Support

- Running an OS requires specific hardware functionality, including:

Execution Modes

Interrupts

Timers and Counters

Atomic Instructions

Privileged Instructions

Virtual Memory

Virtualization Architectures

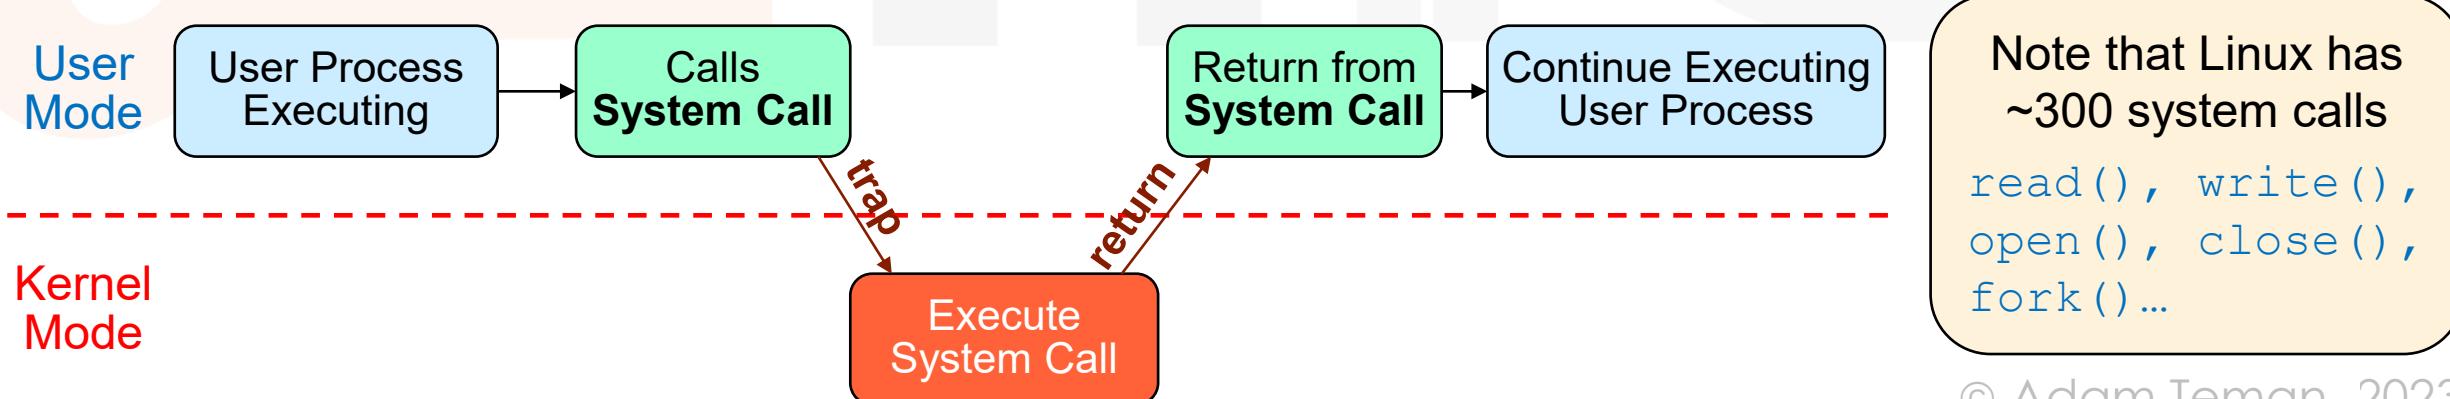
and more...

- RISC-V defines a set of **4096 Control Status Registers (CSRs)** that are primarily used by the OS to support its operations.

- CSRs are accessed by special instructions (e.g., `csrr/csrw` for reading/writing)
- The majority of the CSRs and their functionalities are described in Volume II of the RISC-V Instruction Set Manual, a.k.a., the “**Privileged Spec**”

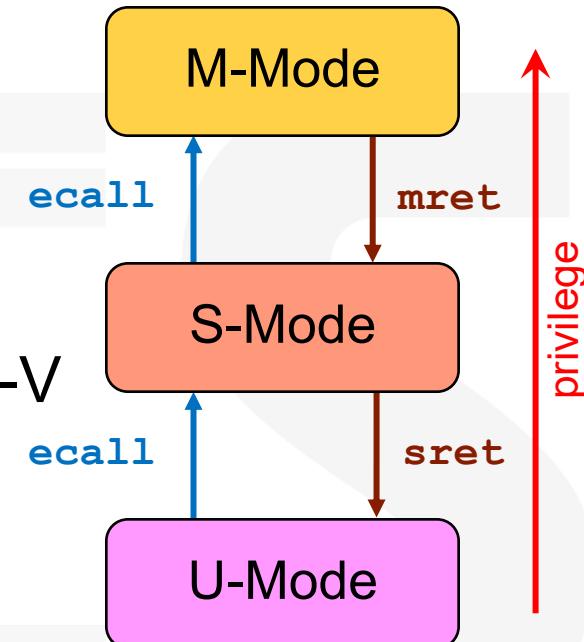
Execution Modes and System Calls

- To provide **privileged operation** to the Kernel, the OS must support at least two distinct **execution modes**:
 - **Kernel Mode**: Only for the kernel, providing **very high access permissions**.
 - **User Mode**: Used by all other applications, allowing **restricted access**.
- A user application is **not allowed** to carry out **privileged operations**
 - Such as reading a file, launching a process, asking for memory (malloc), etc.
 - Therefore, **control must be transferred** to the Kernel running in **Kernel mode**.
- This operation is called a “**System Call**”



Privileged Modes in RISC-V

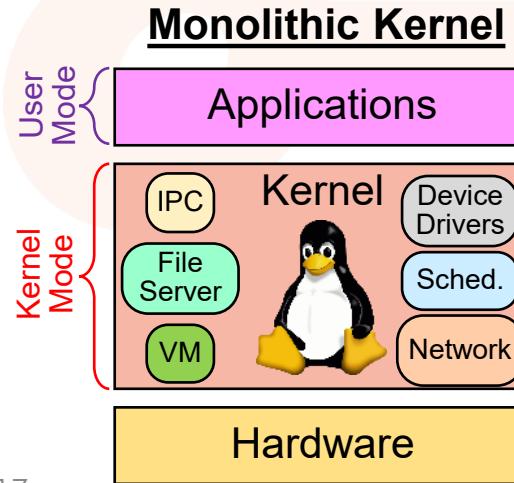
- The RISC-V architecture supports three **execution modes***:
 - User Mode (U) → Lowest Privilege
 - Supervisor Mode (S) → Non-mandatory middle mode
 - Used to abstract away hardware from OS implementation.
 - Machine Mode (M) → Highest Privilege – mandatory in RISC-V
- Each mode has its own set of CSRs, starting with U/S/M
 - For example, `mstatus` is the machine mode status register while `sstatus` is the supervisor mode status register.
- System calls are achieved with the `ecall` and `mret/sret` instructions
 - `ecall` performs a system call from the U/S/M modes to a higher privilege.
 - `mret/sret` return from a system call back to the calling application.



* A fourth mode “Hypervisor Mode” is not yet fully specified. And a special “Debug Mode” is used by debuggers at a high privilege

Kernel Architectures

- There are several options for creating Kernels, trading off **performance, complexity, protection, and more.** The basic architectures are:
 - **Monolithic Kernel:** All main OS operations run as a single program in Kernel Mode, sharing the same address space.
 - **Microkernel:** The Kernel includes only the minimum, essential functionalities of the OS, while other functionalities (services) are run in User Mode.
 - **Hybrid Kernel:** A Microkernel with some additional services inside the Kernel.

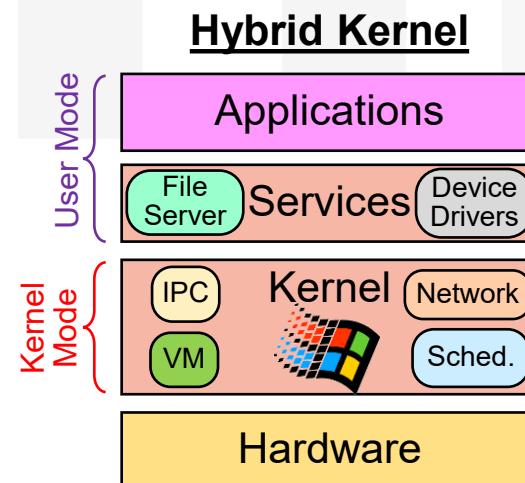


Advantages:

- Fast
- Easy to develop

Disadvantages:

- Large
- Hard to add functionality
- Less Secure
- Failures critical

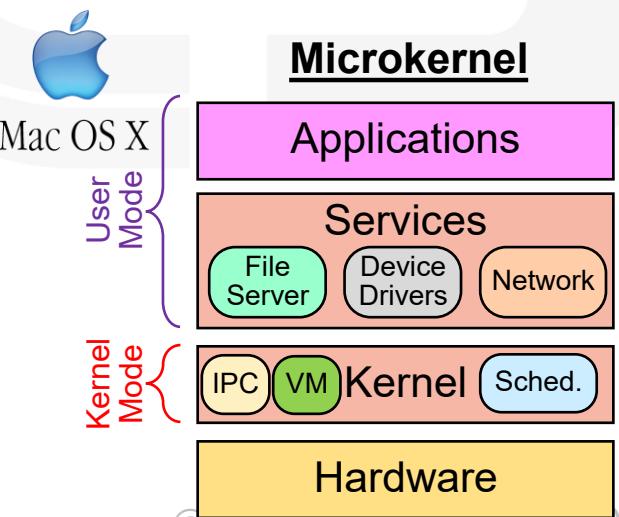


Advantages:

- Small
- Secure
- Modular
- Easy to maintain

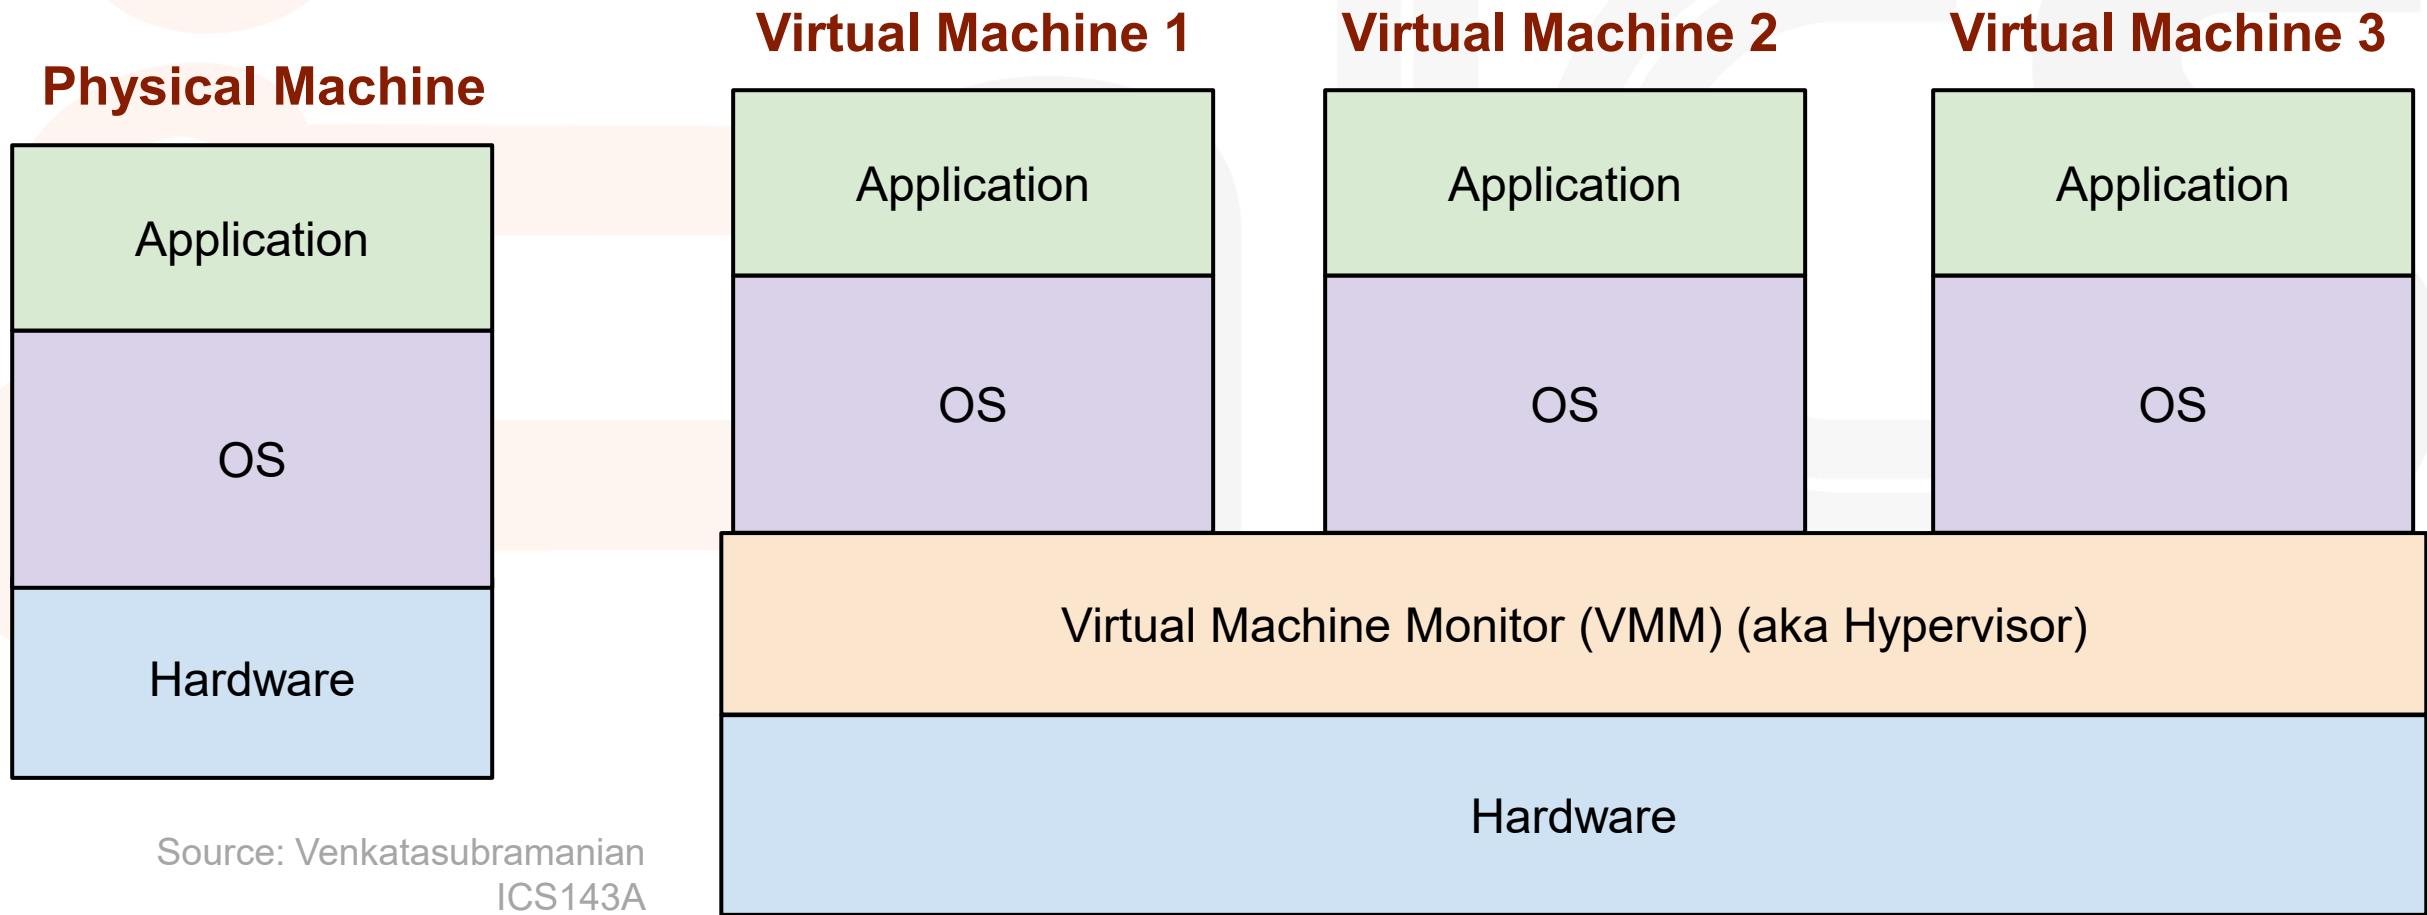
Disadvantages:

- Slow (many system calls)



Virtual Machines

- **Virtual Machines** provide another level of separation from the hardware (*Hypervisor*) to allow running several Operating Systems in parallel.



Introduction

The Kernel

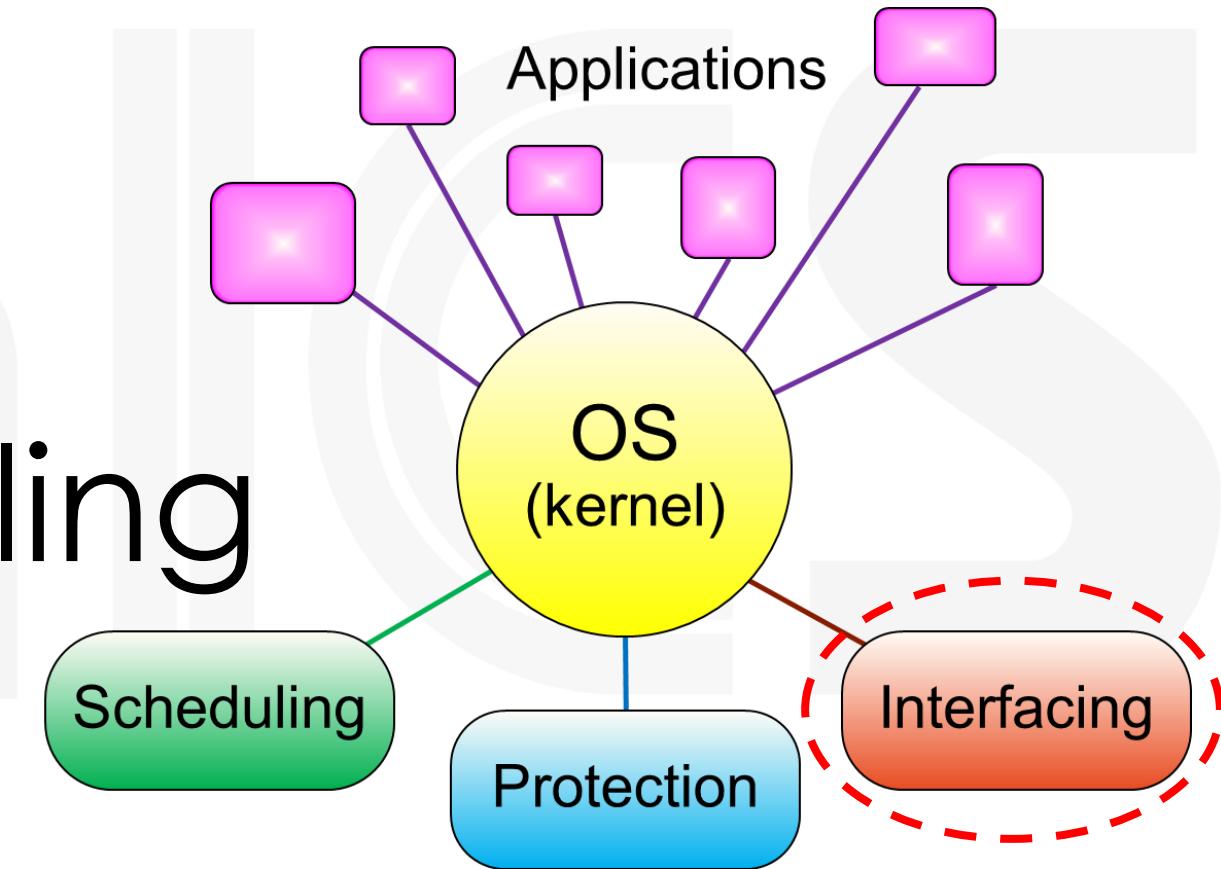
Interrupts

Scheduling

Synchronization

RTOS

Interrupt Handling

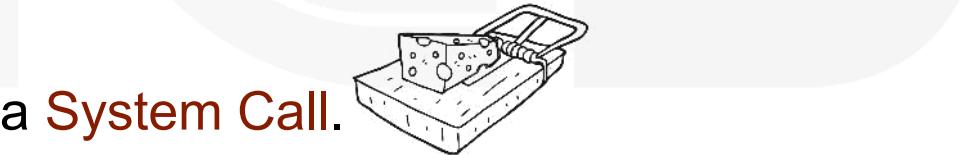


The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



Interrupts and Exceptions

- In the previous section, we were acquainted with **System Calls**
 - A program initiates a **System Call**, when it wants to run a **privileged instruction**.
 - The program's **operation is halted** – *interrupted* – and the **OS** is invoked.
 - This type of interrupt is referred to as a **Software Interrupt** or a **Trap**.
- An **interrupt** is a response by the processor to an event that needs attention from the software.
- We generally differentiate between three categories of interrupts:
 - **Software Interrupts (Traps)**:
 - An **intended exception** thrown by program code, e.g., a **System Call**.
 - **Asynchronous Interrupts**:
 - An interrupt coming from **outside the CPU**, e.g., a **mouse click**.
 - **Synchronous Exceptions**:
 - An **unexpected problem** with the instruction, e.g., **divide-by-zero**.



$$\frac{1}{0} = ?$$

Interrupts: Teman Analogy

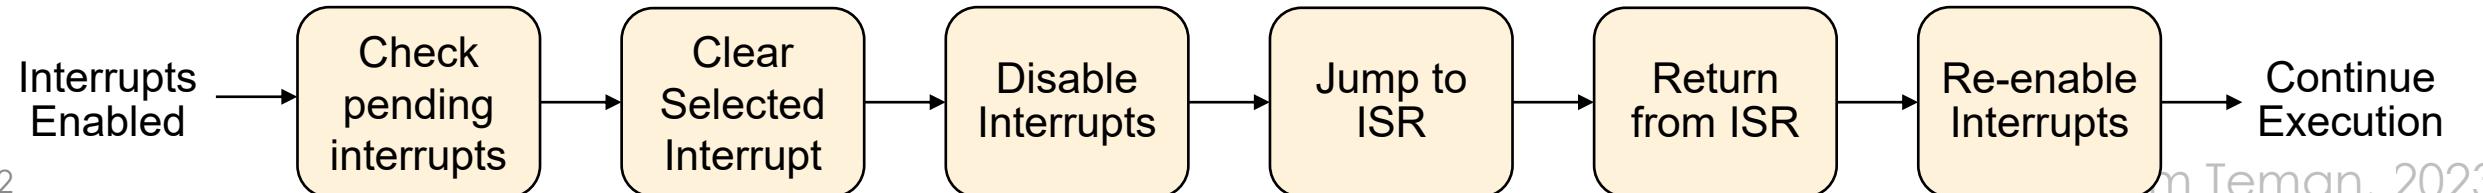
- Ahh... kids are finally asleep...
 - Time for winding down with some TV...
 - Damn! Someone logged out of Netflix...
- System Call!
 - “Honey, what’s the password?”
- Great! A new episode of “Formula 1: Drive to Survive”
 - Wow! Look at Max Verstappen go!
- Exception!
 - Go reset the wifi...
- Okay, Sunday race is starting...
 - Asynchronous Interrupt!
- “Dad, I can’t sleep...”



Interrupt Handling

Interrupts should be as fast as possible, because nothing else can happen when an interrupt is being serviced.

- Upon receiving an interrupt, it must be handled by the operating system:
 - Handling is done by a procedure called an Interrupt Service Routine (ISR).
 - The operating system is invoked (Kernel Mode) and the ISR is called.
 - Synchronous interrupts (traps and exceptions) do this immediately.
 - Asynchronous interrupts go through a slightly different process.
- The operating system decides when to allow asynchronous interrupts
 - If interrupts are enabled, the CPU will check for pending interrupts.
 - If several interrupts are pending, the OS will decide on which one to handle.
 - The selected interrupt is then cleared and interrupts are disabled.
 - Only then, the ISR is called, and upon returning, interrupts are re-enabled.



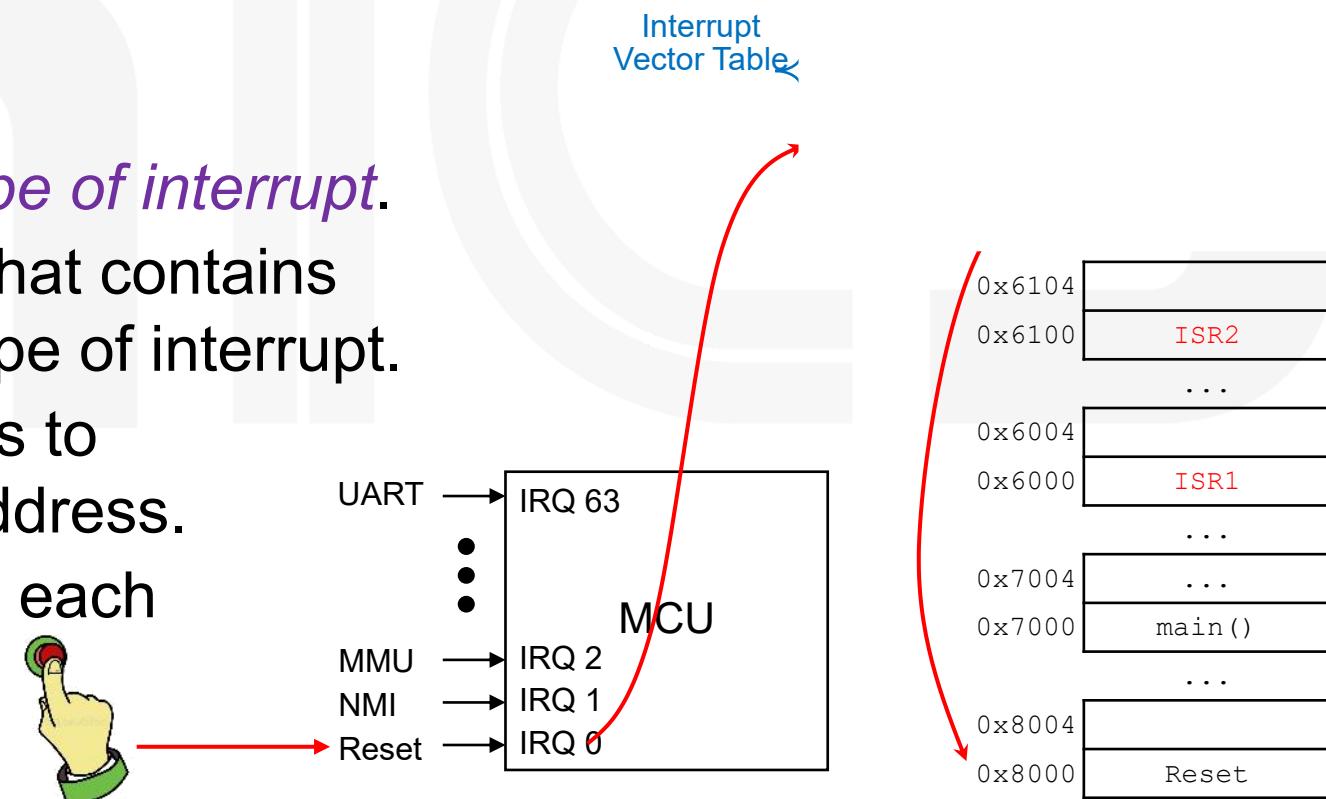
Interrupt Handling – Vectored Interrupts

- Which ISR is called upon receiving an interrupt?
 - Some systems have a single ISR that handles all interrupts.
 - However, this can result in long, complex code → slow interrupt handling.

- The solution...

Indirection, of course!

- Provide a *different ISR* for each *type of interrupt*.
- Use an **interrupt vector** → a table that contains the address of the ISR for each type of interrupt.
- The **encoding of the interrupt** points to the **table entry** that contains this address.
- Usually, this is an address given to each interrupt in a microcontroller.



Interrupt Handling – Priorities

- What happens when several interrupts are pending?

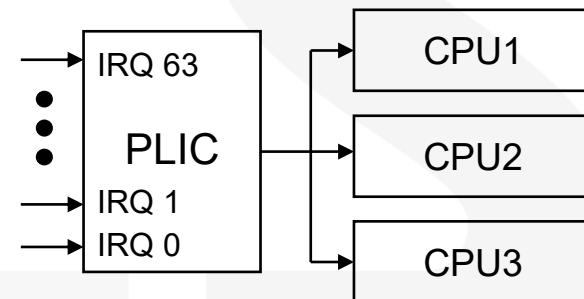
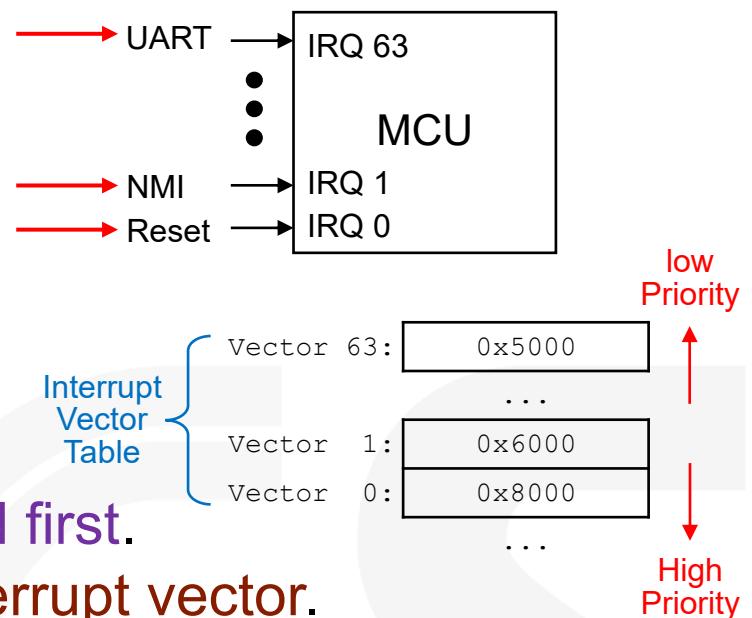
- The system has to decide which interrupt to handle.
- This is often according to priority – the pending interrupt with the highest priority will be handled first.
- In many architectures, the priority is hard-coded into the interrupt vector.

- Is priority enough?

- Not really...
- Interrupt masks may disable certain interrupts and enable others.
- Some interrupts may be deemed non-maskable (NMI).
- If an interrupt arrives while another interrupt is being handled, what should be done?
- What if there are several cores? Several hardware threads (HARTs)?

- Therefore, an interrupt controller is usually used

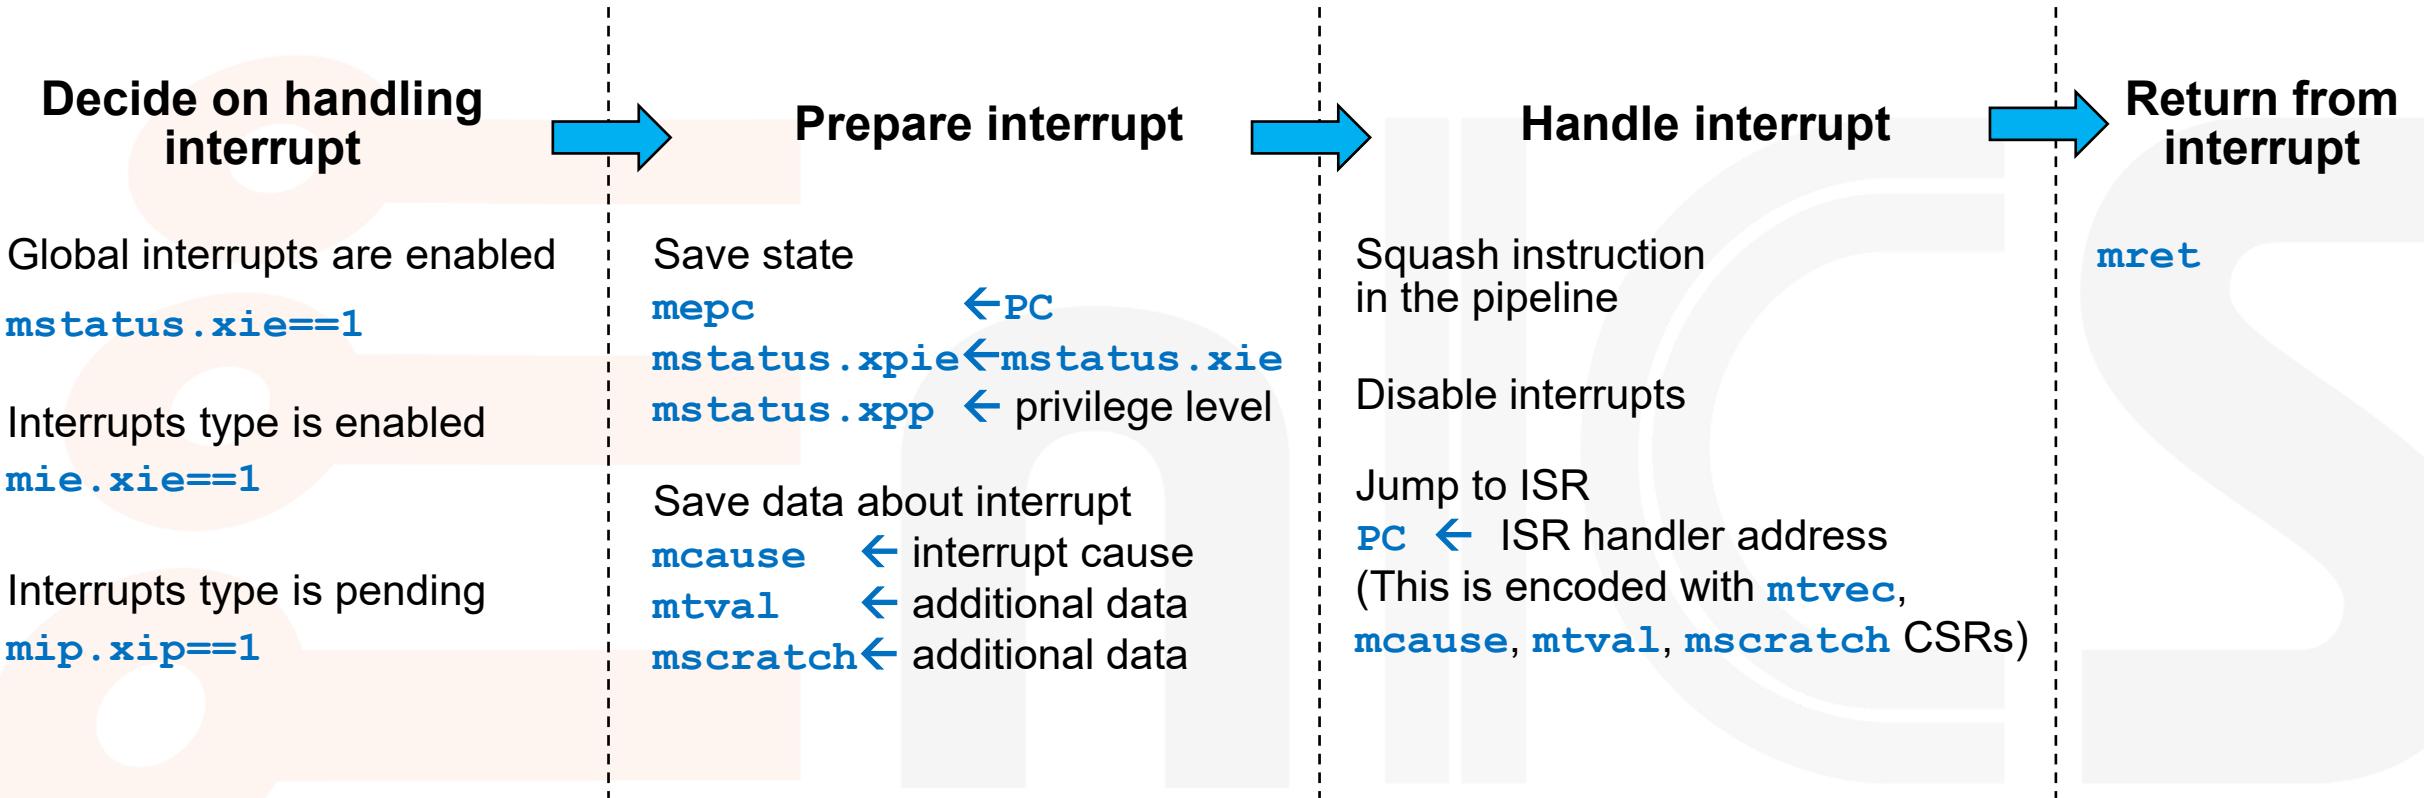
- Usually implemented in hardware, especially for supporting nested interrupts.
- RISC-V specifies a platform-level interrupt controller (PLIC) for multi-core/HART.



Interrupts in RISC-V

- Interrupt handling in RISC-V centers around the **mstatus** CSR
 - Includes bits for interrupt enable (`.xIE`) for each privilege level (`.MIE`, `.SIE`, `.UIE`)
 - Includes bits to store previous state to support nested interrupts:
 - Previous interrupt enable: `xPIE` (`.MPIE`, `.SPIE`, `.UPIE`)
 - Previous privilege level: `xPP` (`.SPP`, `.MPP`)
 - Return from interrupt instructions (`mret`, `sret`) restore state upon return.
- Pending interrupts are listed in the **mip** CSR.
 - Separated into external (`.meip`), software (`.msip`), and timer (`.mtip`) bits.
 - Corresponding enable bits (`.meie`, `.msie`, `.mtie`) in the **mei** CSR.
- Additional CSRs store information about the interrupt
 - **mcause**: cause of the interrupt, **mtvec**: address of the interrupt handler
 - **mepc**: address of interrupted instruction, **mtval**, **mscratch**: more info

Interrupts in RISC-V



Introduction

The Kernel

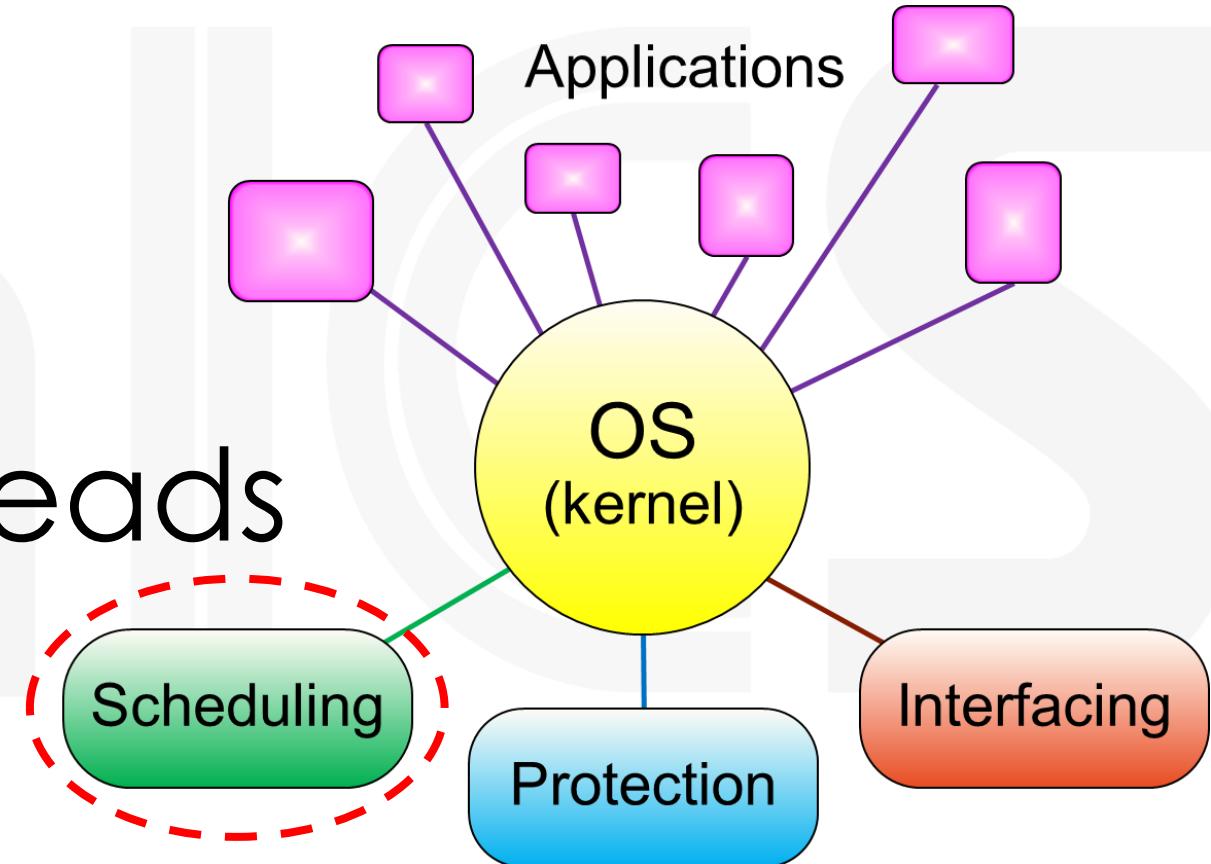
Interrupts

Scheduling

Synchronization

RTOS

Scheduling: Processes and Threads



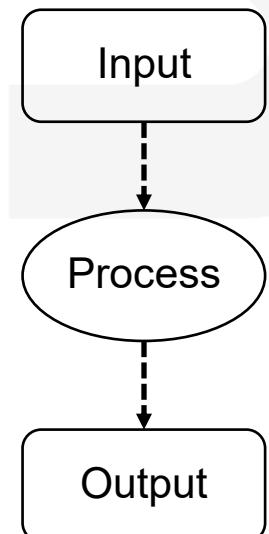
The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



Concurrency

- Many things are going on in an Operating System
 - Applications, interrupts, background tasks, maintenance...
- However, the CPU can only run one program at a time
- Solution #1: Cyclic Executive
 - An endless loop that switches between tasks
 - This is really hard to write, maintain and optimize.
 - Can't handle interrupts!
- Solution #2: Process Abstraction
 - Decompose complex problems into simple ones (programs).
 - Create an instance of the program ("a process") and run it.
 - Timeshare between running processes through scheduling.
 - Each process feels like it has the entire CPU.

```
while true {  
    do part of task 1;  
    do part of task 2;  
    do part of task 3;  
}
```



Programs and Processes

- A **program** is just the code
 - A binary, including the **text** and **static data**.
 - Totally passive. Just bytes on a disk.
- A **process** is an instance of a program in execution
 - A copy of the **program code** (**text** and **static data**)
 - The **processor state** (**PC**, **registers**, **memory**)
 - **Resources** (**file pointers**, etc.)
- A **program becomes a process when an executable is loaded into memory**
 - There may be many processes running copies of the same program.
 - e.g., each tab of Google Chrome is a separate process
- However, there is **only one set of resources**
 - Process must be *scheduled* and must ensure *protection*.

Program

```
main () {  
    ...  
    foo();  
    ...  
    bar();  
}
```

Process

```
main () {  
    ...  
    foo();  
    ...  
    bar();  
}
```

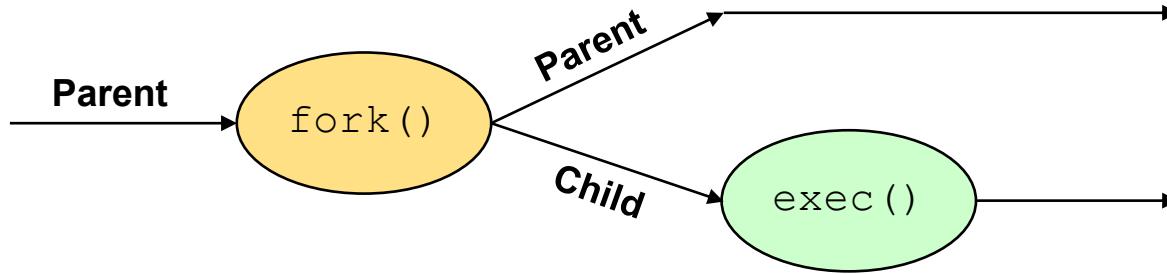
Registers

Address Space

Resources

Name	11% CPU	37% Memory
Apps (6)		
Google Chrome (13)	0.1%	1,620.1 MB
Google Chrome	0%	30.7 MB
Google Chrome	0.1%	145.3 MB
Google Chrome	0%	156.7 MB

Process API



- **How is a process created?**

- In Posix systems, the most common way to create a process is to `fork` it.
- `fork()` is a system call that **creates a copy** of process.
- The new **child process** is an almost exact copy.

- **Why do we want a copy of the parent process?**

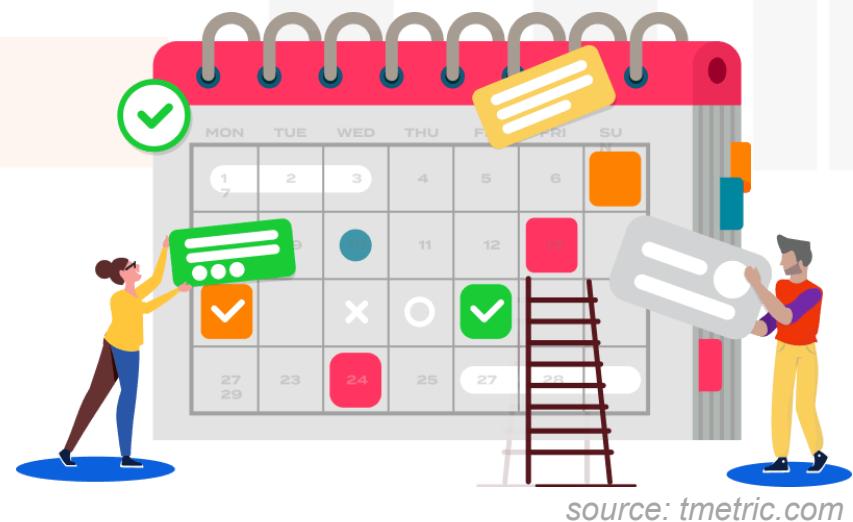
- Actually... we don't.
- The `fork()` call is followed by an `exec()` call, which **runs a new program**.

- **Why create a copy of the parent process?**

- **Very good question...** without a really good answer.
- The short answer... because this API just works (“[Lampson’s Law](#)”)
- The better answer... because it lets us **run things on a child process before executing it**. This is what enables “pipes” (`|`) and “redirects” (`>`) in Unix.

Scheduling

- A CPU can only run **one process at a time**
 - How does the OS **choose** which process will be executing?
 - When should the OS **change** the currently running process?
 - How can the OS efficiently **utilize** the resources?
- The decision of **which task to run and when** is called **Scheduling**
 - A **scheduling policy** tries to optimize **performance** and **fairness**.



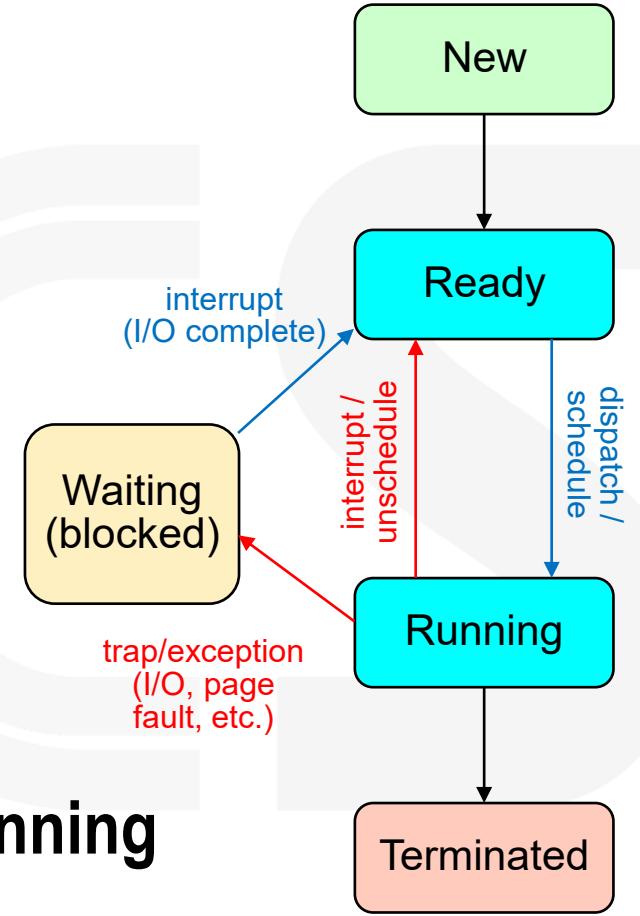
Scheduling Policies

- The scheduling policy dictates which process will be run next.
- This may depend on what we want to achieve...
 - Maximum throughput, minimize response time, fairness, etc.
- Therefore, various scheduling policies exist:
 - First-come First-served (FCFS)
 - Shortest Job First (SJF)
 - Shortest Remaining Processing Time First (SRPT)
 - Round Robin (RR)
 - Priority
- The most common approach today is Multi-level Feedback Queues (MLFQ)
 - Provide several queues with different priorities
 - Apply RR within the queues. Move between priorities based on aging.



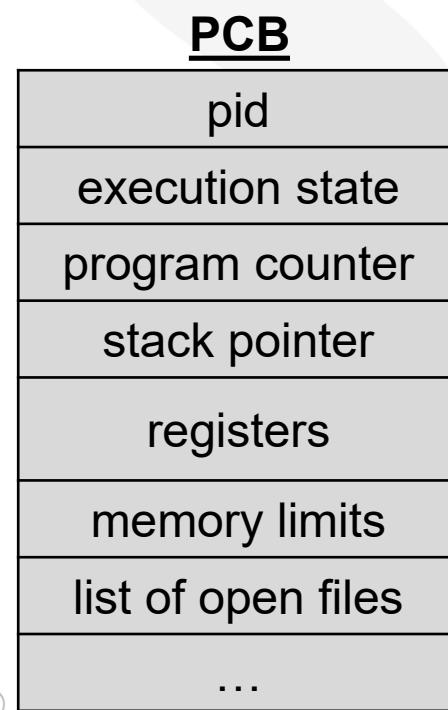
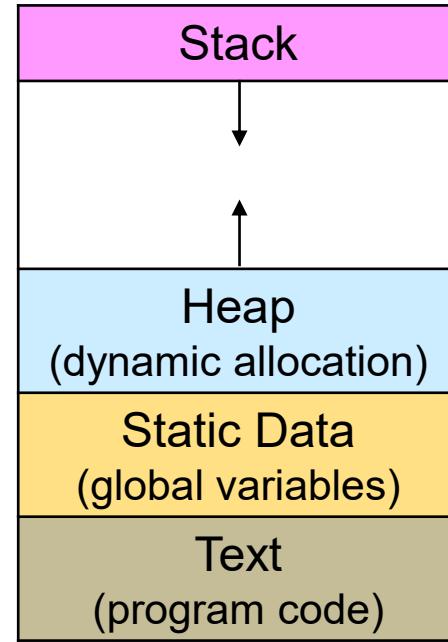
Queue-based preemptive scheduling

- Some policies feature **queues of processes waiting to run**
 - Queues for each device, timers, interrupts, messages...
- Each process has an **execution state** within the queue:
 - **Ready**: Waiting to be assigned to a CPU
 - Could run, but another process has the CPU
 - **Running**: Executing on a CPU
 - It's the process that currently controls the CPU
 - **Waiting (a.k.a., “blocked”)**: waiting for an event, e.g., I/O completion, or a message from another process
- The OS periodically decides to change which process is running
 - This is known as “**preemptive scheduling**”:
 - “**Preempt**”: *to acquire or appropriate before someone else.* ([dictionary.com](https://www.dictionary.com))
- Preemptive scheduling requires a mechanism called a “**Context Switch**”



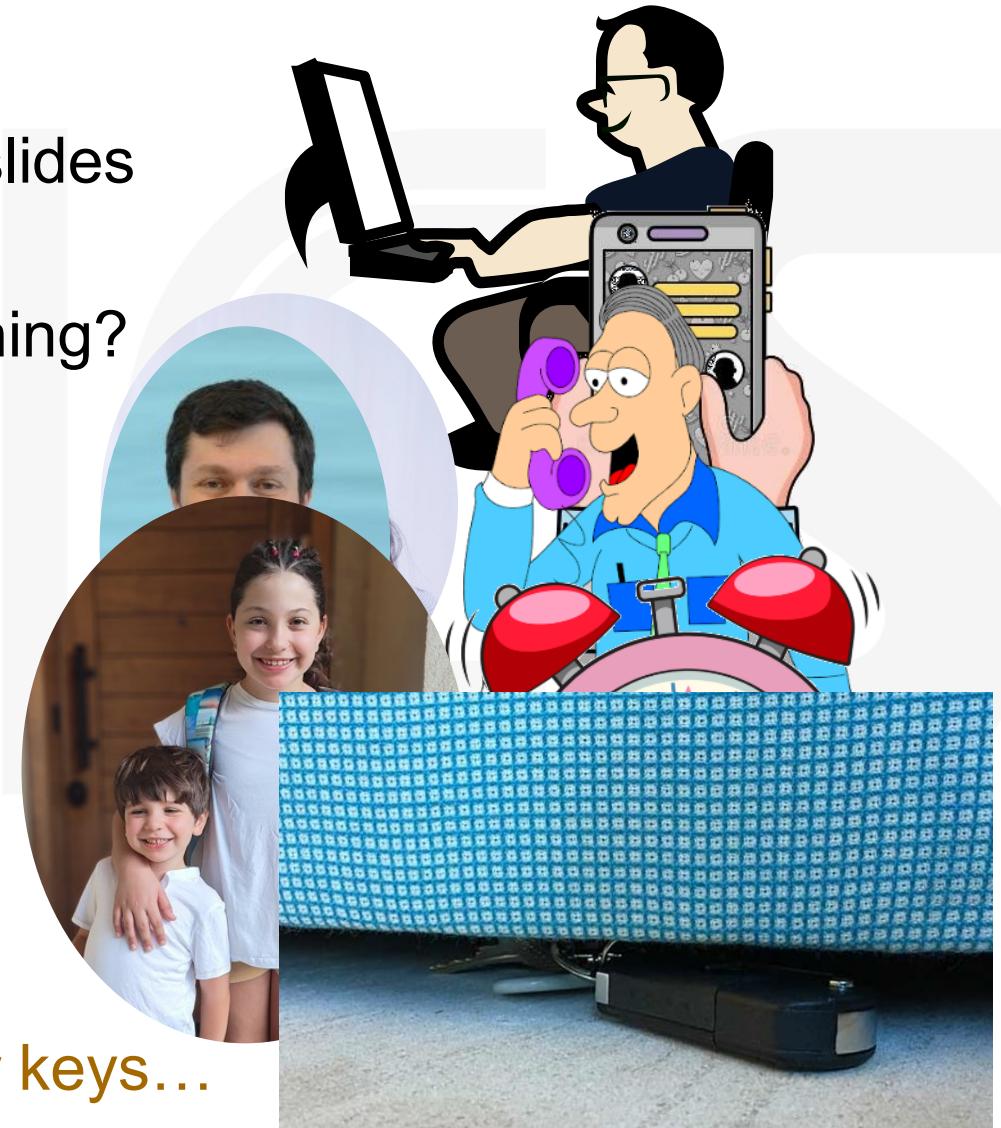
Context Switching and the PCB

- **Context Switching** is the mechanism to change which process is currently running on the CPU
 - Invoke the kernel, store the state of the running process, load the new process onto the CPU.
- Everything necessary for loading/storing a process is stored in a data structure called a **process control block (PCB)**, including:
 - The code for running the program (“text”)
 - The CPU registers (e.g., PC, SP, etc.)
 - The address space (static data, stack, heap, page tables)
 - I/O Status (open files, I/O devices allocated to the process)
- And additional information used by the OS.
 - Process identification (PID)
 - Process execution state (ready, running, waiting, ...)
 - Scheduling information (priorities, queue pointers, ...)
 - Accounting data (e.g., CPU used, running time, ...)



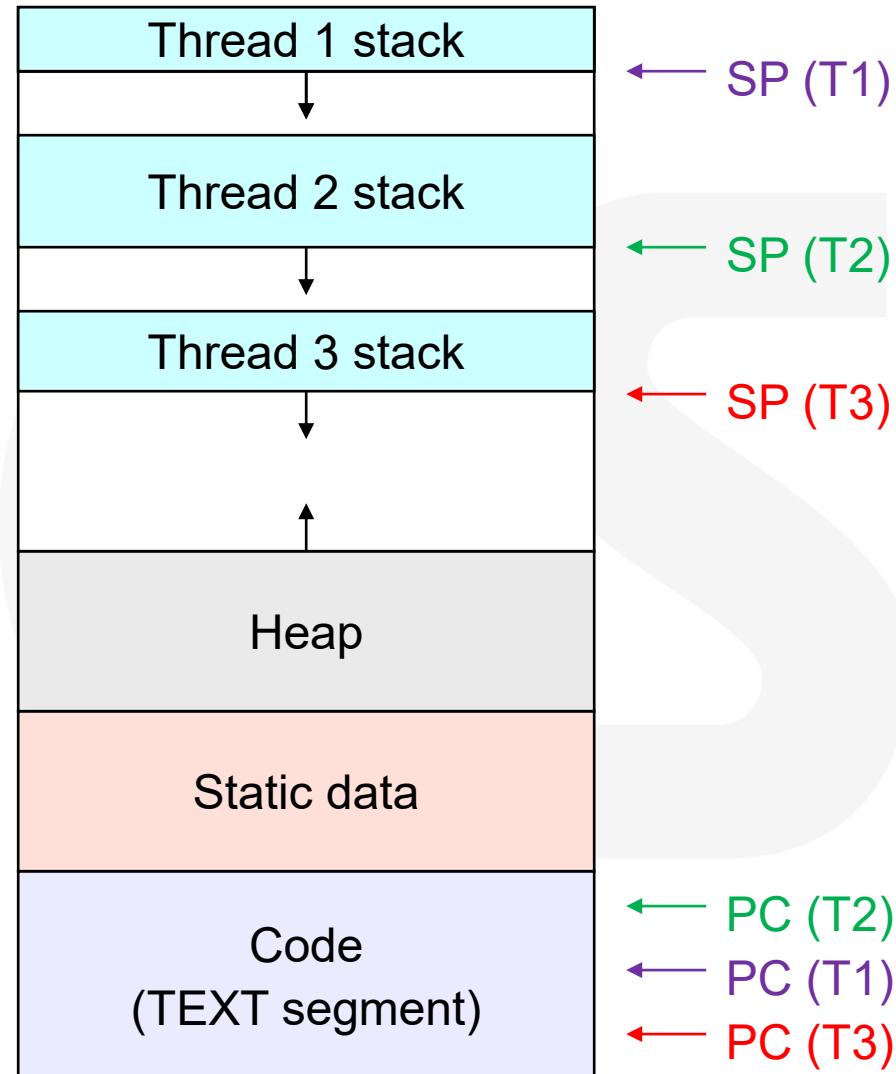
Context Switch: Back to the Teman Analogy

- I need to prepare this lecture...
 - Open up PowerPoint... Start working on the slides
- WhatsApp message
 - It's the wife... She's at the store... need anything?
 - Save work... go to check in the kitchen...
 - Back to work... where was I?
- Phone rings...
 - It's Alex... need to send that email – now!
 - Save work... open mail client... send mail
 - Okay, re-open PowerPoint... continue editing
- Alarm clock rings...
 - Time to pick up kids...
 - Save work... put computer to sleep... find car keys...



The Thread

- Processes do not share resources well
 - Context switching is expensive
 - e.g., saving and restoring a page table
 - Sharing data between processes is complicated
- Idea: Separate **concurrency** from **protection**
 - Create a *lightweight process* → a “**Thread**”
 - A **process** can have several **threads**
- A **process** has a **single address space**
 - Code and Data sections
 - OS resources (open files, signals)
- Each **thread** has its own **PC, CPU registers, stack**
 - There is no protection between **threads**



The Thread (cont.)

- **Threads** are concurrent executions sharing an address space

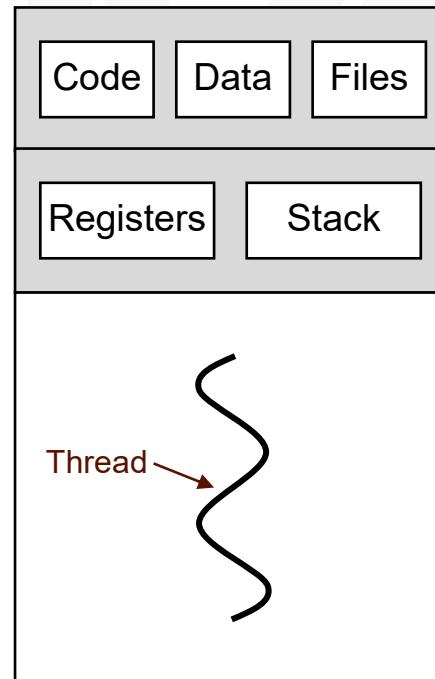
(and some OS resources)

- **Threads** are *cheaper than processes*:

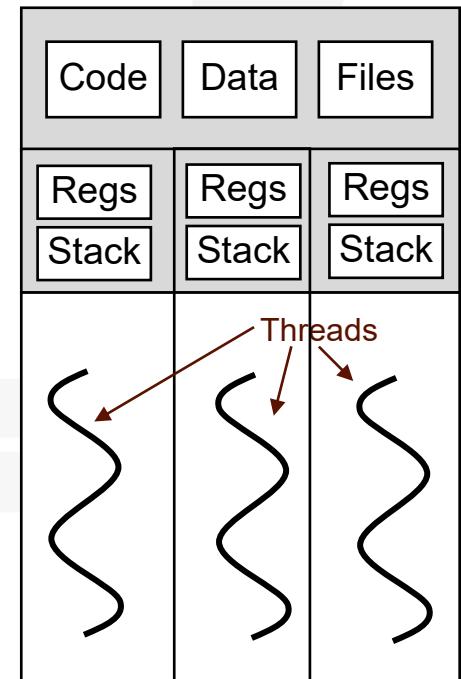
- Creating **threads** is cheap
- Context switching **threads** is cheap
- Sharing data between **threads** is cheap

- In a **multiple threaded task**, while one **thread is blocked and waiting**, a second **thread in the same process can run**.

- Higher throughput
- Improved performance



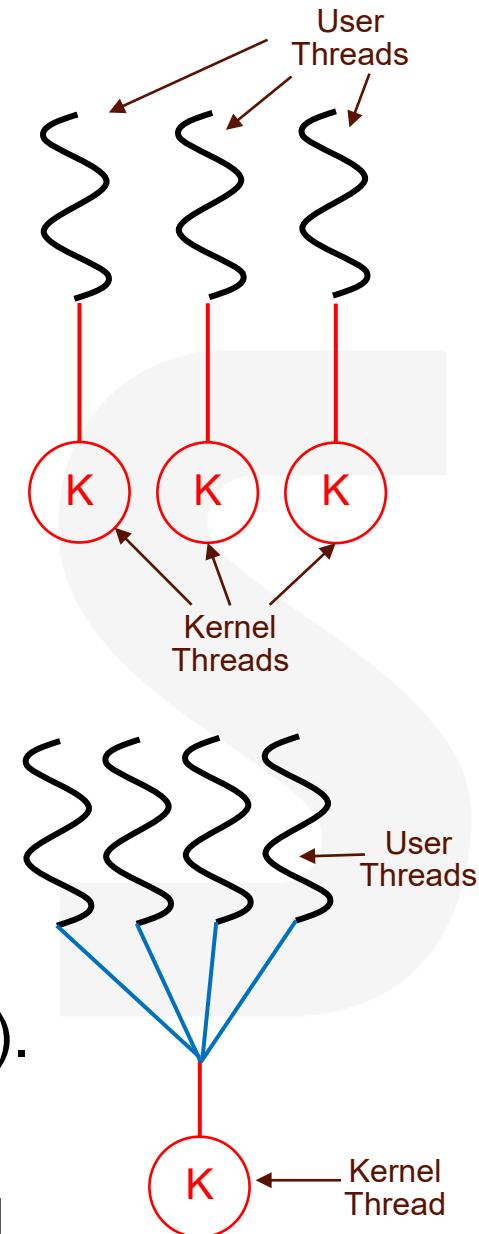
Single-Threaded
Process



Multi-Threaded
Process

Kernel vs. User-Level Threads

- How are threads managed?
 - How are they **created**? Who allocates memory for them?
 - How are they **context switched**?
- In a **Kernel Thread**, the OS manages the threads
 - Each thread has a kernel thread associated with it.
 - All thread operations are implemented **in the kernel**.
 - Thread operations are **cheaper** than process ones, but still **costly**.
- Alternatively, a process can manage its threads in **user space**
 - **User-level threads** are much cheaper to manage (**10-100x faster!**).
 - Can be limited due to I/O operations going through the OS.
- **Hardware threads (HARTs)** provide very efficient context switching
 - Also called **simultaneous multi-threading** (SMT) or **hyperthreading** (Intel)



Introduction

The Kernel

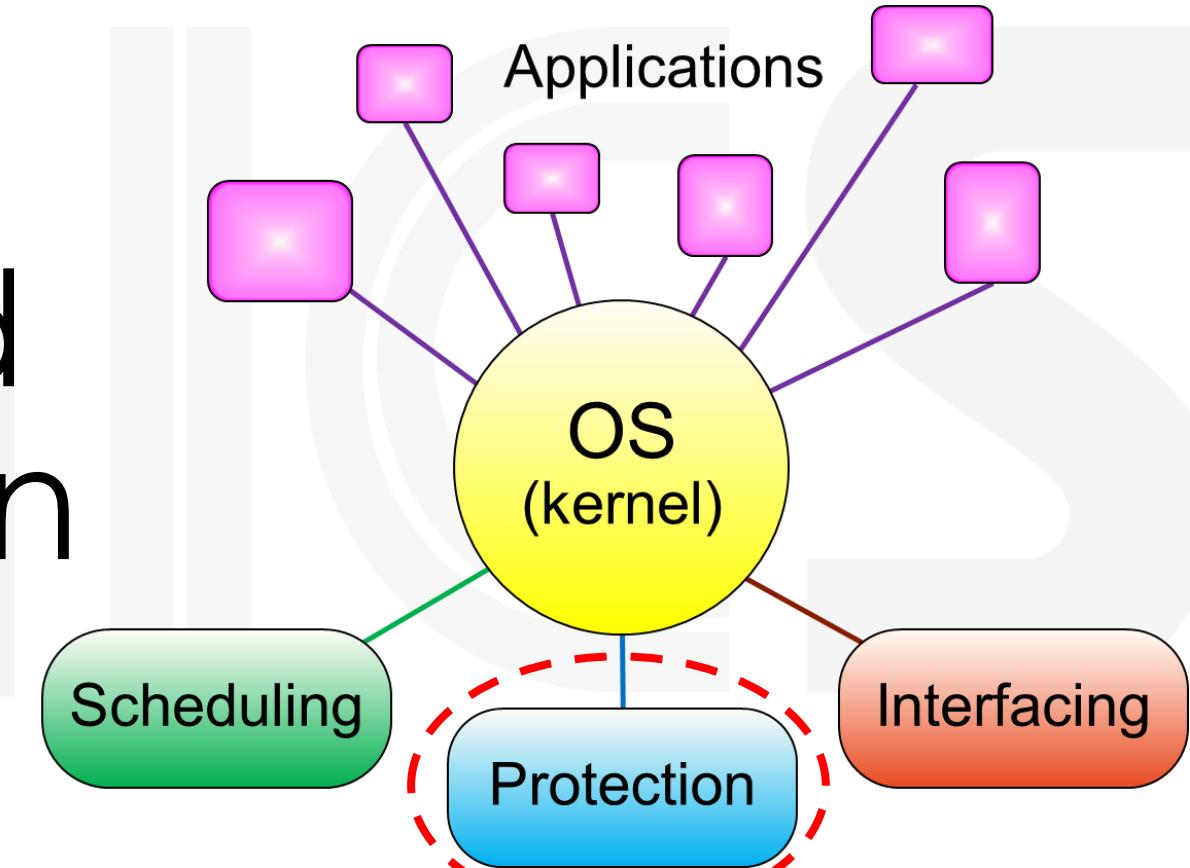
Interrupts

Scheduling

Synchronization

RTOS

Protection and Synchronization



The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



Protection

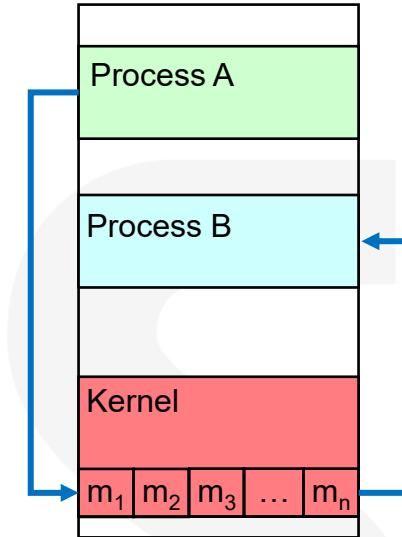


- Protection is required for:
 - To prevent errant or malicious misuse of the system by users or programs
 - To ensure that shared resources are used according to system policies
 - To limit the damage caused by errant programs
- The basic means of protection is
 - to provide each process with its own separate address space
 - Only the kernel (supervisor) has access to the full physical address space
 - The process can only address the space defined in its PCB
 - Virtual memory and page tables help abstract away the physical memory from the program
- But how can we share data between processes and maintain protection?
 - In other words, how do we enable inter-process communication (IPC)?

Inter-Process Communication Models

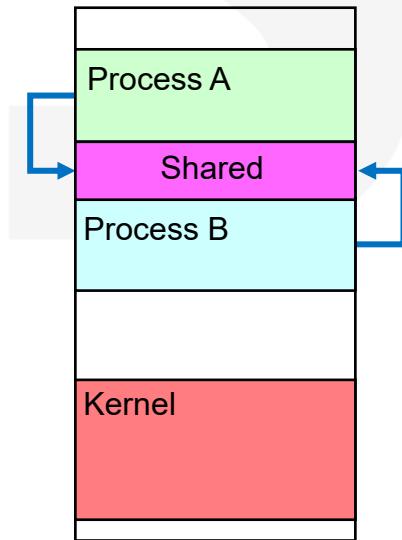
- **Option #1: The Message Passing model**

- **Let the Kernel handle it!**
 - `send()` – store data on a queue managed by the kernel.
 - `receive()` – read data from the queue through the kernel.
 - But this suffers from **a lot of overhead** due to frequent system calls.



- **Option #2: The Shared Memory model**

- **Let several processes access the same memory space.**
 - No kernel intervention is required.
 - But how can we **ensure protection** on the shared memory?



Can we use a queue?

- In a **Producer-Consumer** model, we can restrict write access

- Only the producer can write to the queue.
- For example, use a First-in First-out (FIFO) buffer.
- This enables passing messages between tasks without accidentally corrupting each other's messages.

- This works for two processes

- But what if there are **several producers** and/or **several consumers**?
- They need to **share** the FIFO pointers.
- What will happen if **both try to update** the pointer at the same time?

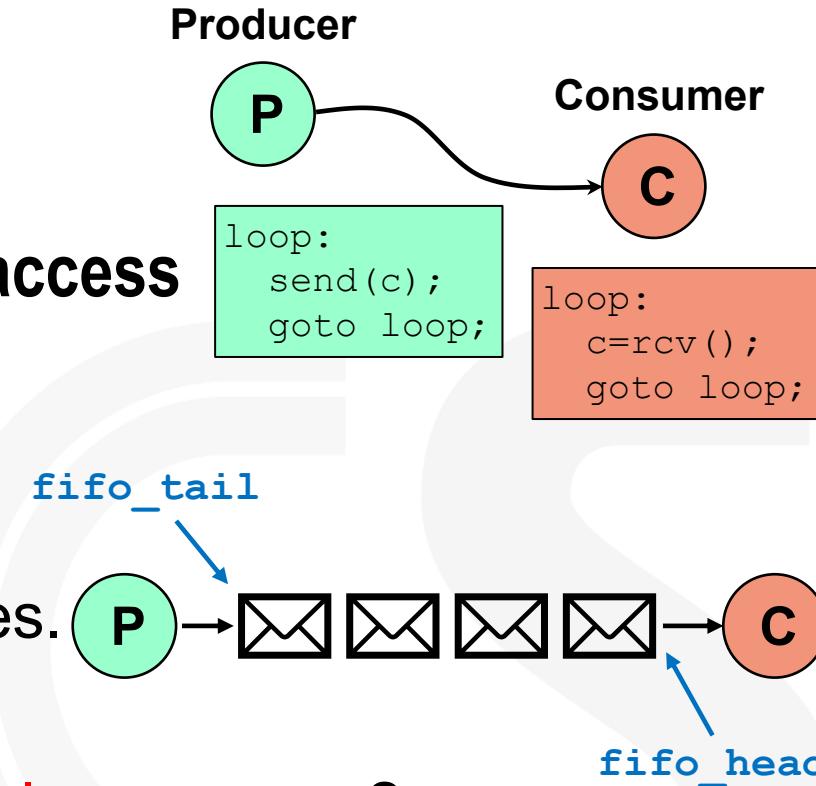
Producer 1

```
...
//Producer 1 pushes to FIFO
*fifo_tail = producer1_data;
fifo_tail++;
...
```

What happens if we context switch here?

Producer 2

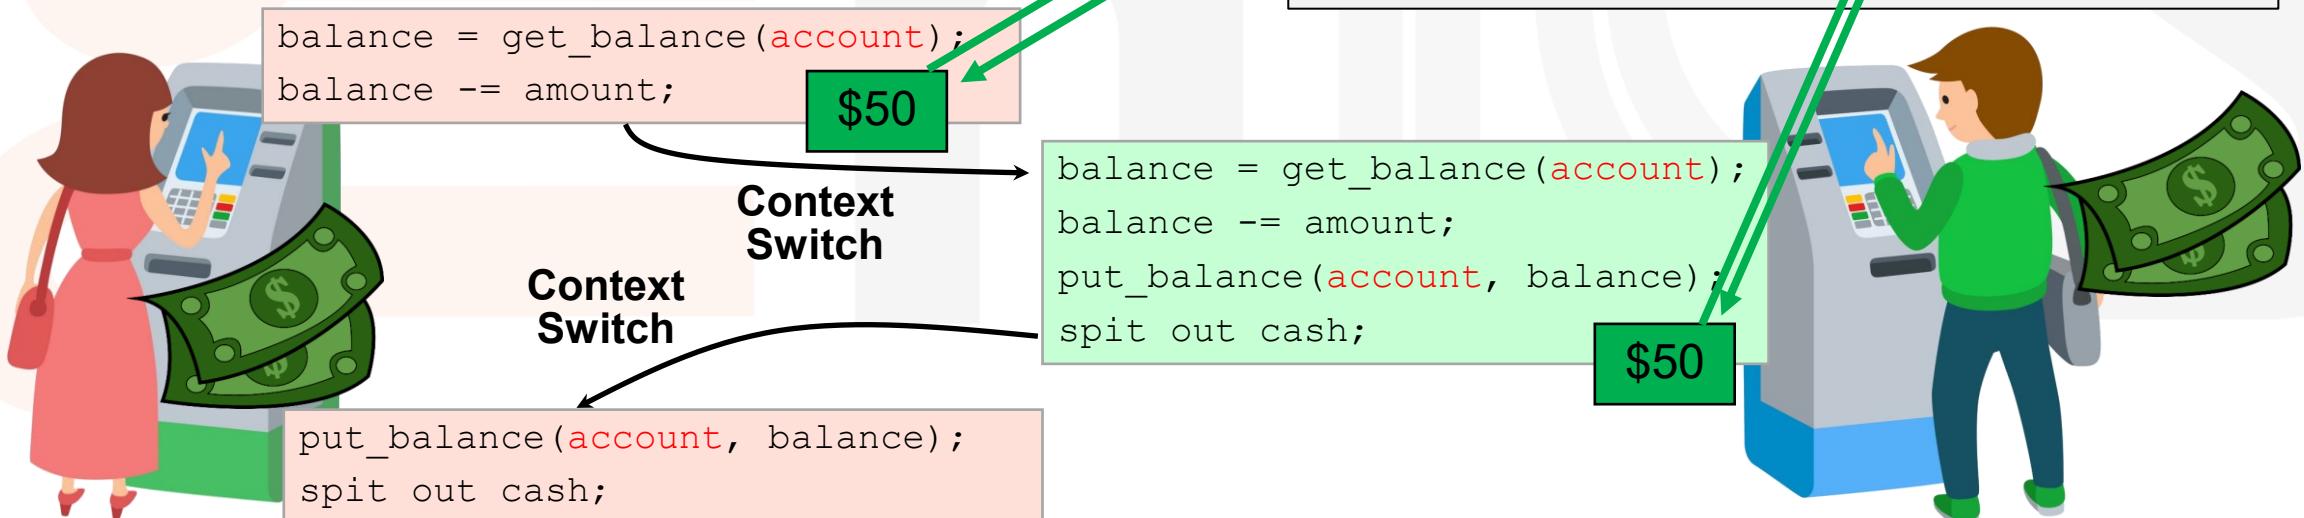
```
...
//Producer 2 pushes to FIFO
*fifo_tail = producer2_data;
fifo_tail++;
...
```



The Critical Section

- Simultaneous access to a shared variable can result in a **race condition**.

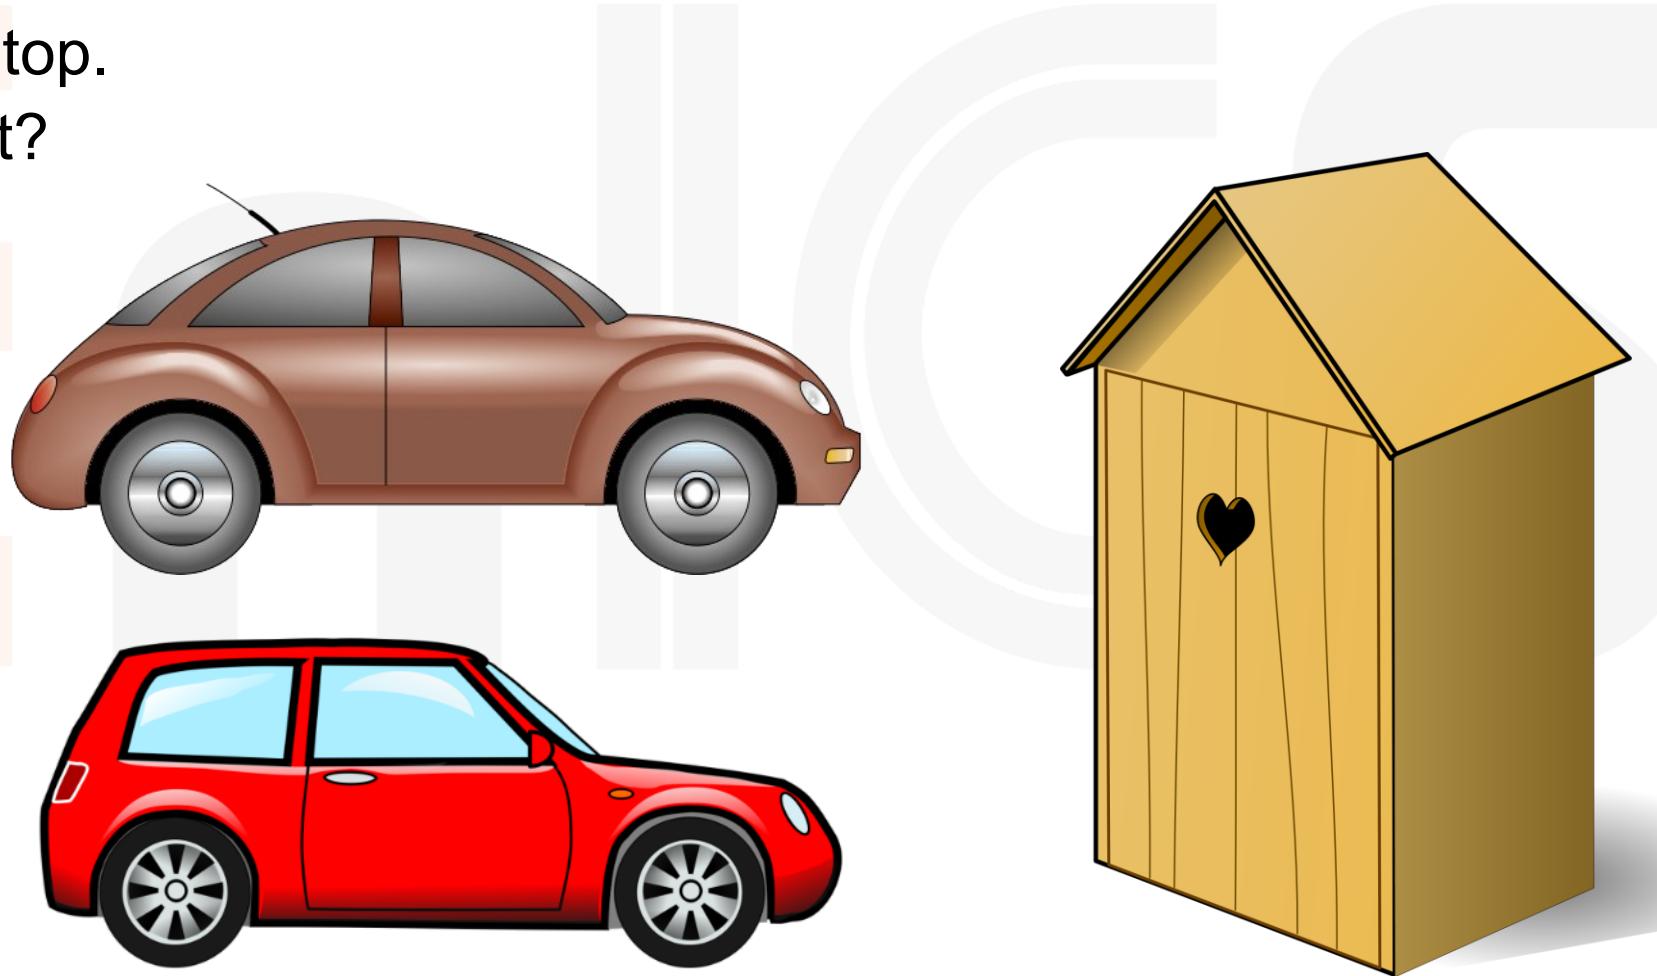
- For example, the wife and I both withdraw \$50 from the ATM.
- Preemptive scheduling could lead to the following course of action:



- The code block where a shared variable is accessed is called “**The Critical Section**”

Gas Station Analogy: Critical Section

- A gas station on the highway has a public bathroom.
 - Two cars make a pit stop.
 - Who will get there first?
- The bathroom is in the critical section!



The Critical Section Problem

- A **Race Condition** occurs if:
 - Several processes/threads manipulate a **shared resource** concurrently, and
 - The outcome depends on the **particular order** in which the access takes place.
- **Synchronization** around critical sections is needed to prevent race conditions.
- The criteria to solve the critical section problem are:
 - **Mutual Exclusion**: At most one thread can be inside its critical section at any given time.
 - **Progress**: A thread outside its critical section, cannot prevent another thread from entering the critical section.
 - **Bounded Waiting**: A thread waiting on its critical section will eventually be allowed to enter it.

```
do {  
    :  
    Entry Section  
    Critical Section  
    Exit Section  
    :  
} while {1};
```

Mutual Exclusion (Mutex)

- A Mutex is a primitive that enables synchronized access to a shared resource (i.e., variable).

- Define a lock:

- `acquire()`: obtain the right to enter the critical section
- `release()`: give up the right to be in the critical section

- Often accessed with a “**Spinlock**”

- Busy-wait on the lock until it is free.

- Problem: Spinlocks have critical sections, too!

- `acquire()`/`release()` must be **atomic**
- **Atomic** == executed without interruption

- Need help from the hardware

- *Atomic instructions*, e.g., `test-and-set`
- or `disable interrupts` to prevent context switches

What happens if we context switch here?

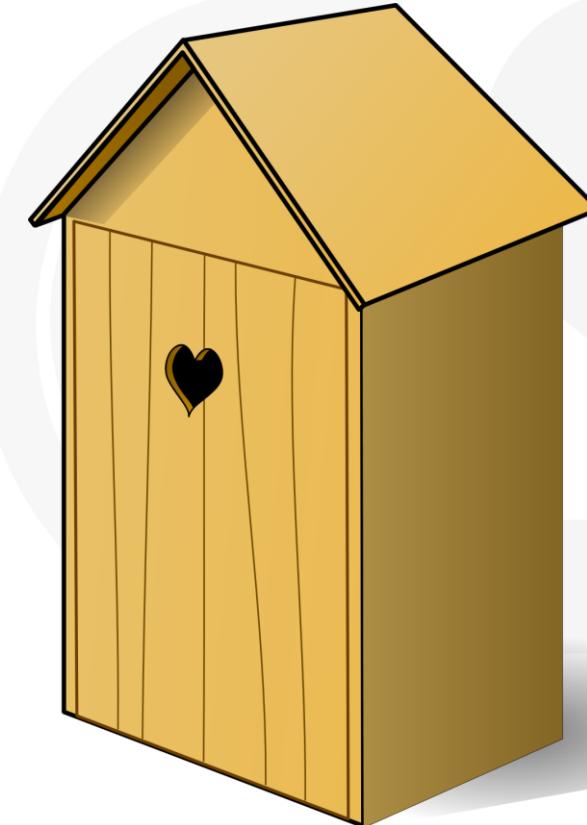
```
struct lock {  
    int held = 0; }  
  
void acquire(lock) {  
    while (lock->held);  
    lock->held = 1; }  
  
void release(lock) {  
    lock->held = 0; }
```

```
bool test_and_set(bool *flag) {  
    bool old = *flag;  
    *flag = True;  
    return old; }  
  
struct lock {  
    int held = 0; }  
  
void acquire(lock) {  
    while(test_and_set(&lock->held));  
}  
  
void release(lock) {  
    lock->held = 0; }
```

Gas Station Analogy: Mutex

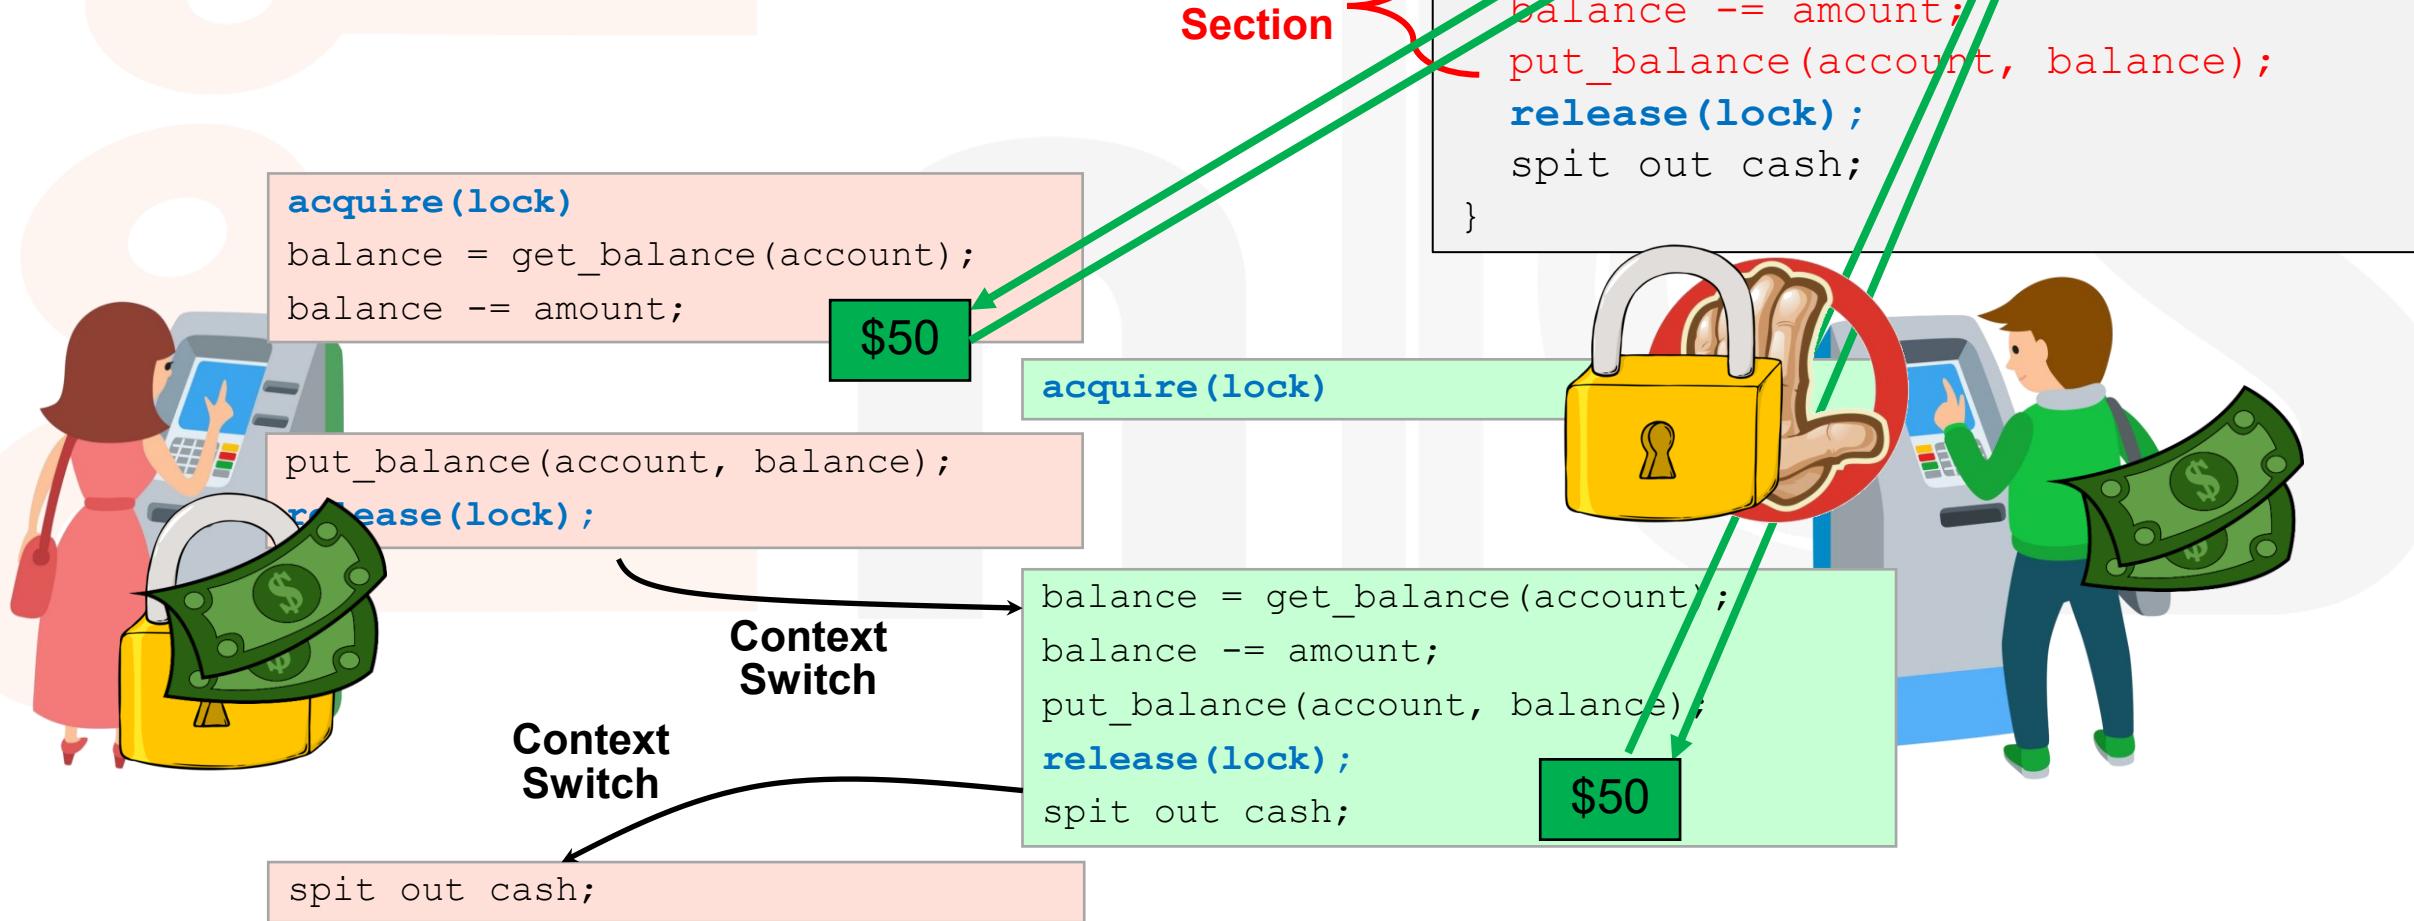
- Gas station has one bathroom and one key.

- First customer acquires key.
- Second customer waits.
- Until first customer releases the key.
- Now second customer can acquire the key.



Let's go back to our ATM

- We'll protect our critical section with a lock.

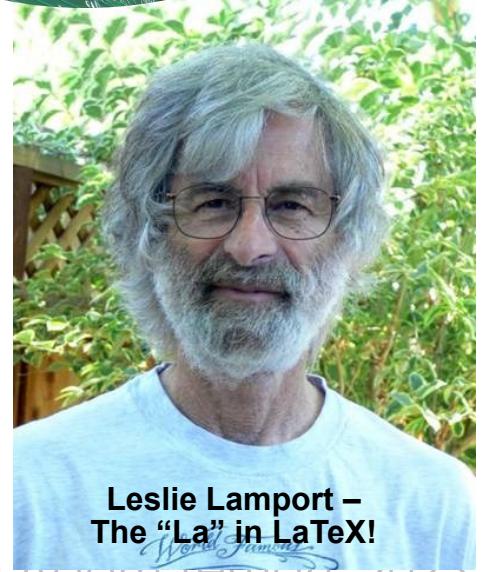


Implementing Synchronization

- **Synchronization** is a true “black hole” that should not be approached by the faint-hearted.
 - Carefully written “[synchronization libraries](#)” enable programmers to implement mutual exclusion and write race-free code.
- Mutual exclusion can be achieved through software only*
 - For example, [Lamport’s Bakery Algorithm](#) (1974)
 - But these algorithms require many processor cycles to implement.
- Therefore, locks are usually implemented using “[atomic operations](#)”
 - Sequences that are guaranteed to complete within a single indivisible operation.
- OS classes just “assume” the existence of such instructions...
 - But how do we actually implement them in hardware?



Source:
wikipedia



Leslie Lamport –
The “La” in LaTeX!

Atomic Instructions

Three main types of atomic instructions are found in synchronization libraries:

- **Atomic Exchange:**

- Swap the content of a register and a memory location.
- If the lock was free – it returns a ‘0’, otherwise we **spin**.
- But this means we continually write to memory, **which is very bad...**

Acquiring a lock with
Atomic Exchange

```
R1=1;  
while (R1==1)  
    EXCH R1, lock;
```

- **Test and Set:**

- Read the lock, and if it is free, acquire it.
- No need to write to memory while spinning.
- But requires **load and store within single instruction**.
- Would lead to deep pipeline / performance hit.

Test and Set Command

```
if (*lock==0) {  
    *lock=1;  
    R1=0;  
} else {  
    R1=1;  
}
```

- **Load Linked / Store Conditional (LL/SC):**

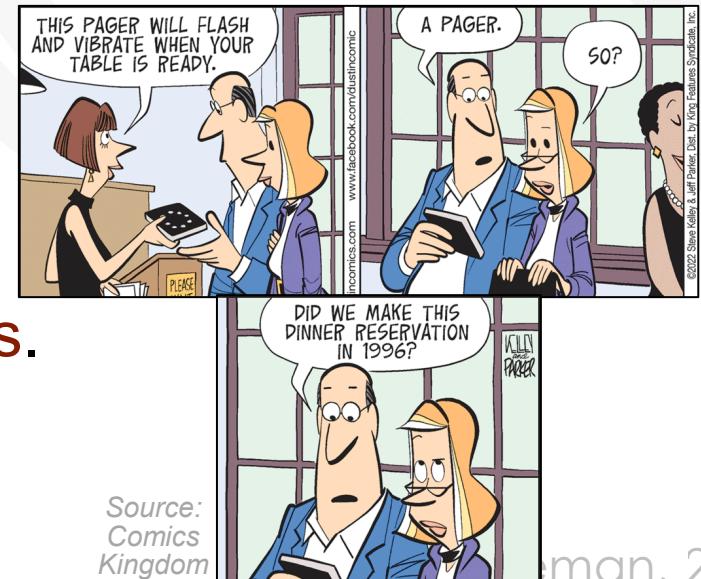
- Split the operation into **two commands** to simplify memory access.
- Let’s see how this works...

Load Linked / Store Conditional

- The LL/SC approach creates an atomic operation with two instructions.
- Load Linked (or Load Reserved)
 - Carry out regular load from memory to a register.
 - Create a “reservation” on the load.
- Store Conditional
 - Attempt to store a value to the memory location.
 - If the “reservation was broken”, the store fails.
- What “breaks a reservation”?
 - A context switch (→breaks atomicity)
 - A store (by another processor) to the reserved address.
- The reservation is held in a hidden register and can be accessed in parallel to the memory access.

Implementing a lock with LL/LR

```
void lock (&lockvar) {  
    trylock:  
        MOV R1,1           // R1=1  
        LL   R2, lockvar  // R2=*lockvar  
        BNE R2, trylock   // spin if R2!=0  
        SC   R1, lockvar  // try *lockvar=1  
        BEQ R1, trylock   // spin if failed
```



Atomic Operations in RISC-V

- RISC-V adopted the LL/SC approach

- Load-Reserved (`lr`):

Reservation includes address of memory access

- Store-Conditional (`sc`):

Releases reservation and returns 0 if successful.

```
lr.w rd, (rs1) # rd= Mem[rs1]
```

```
sc.w rd, rs2, (rs1) # Mem[rs1] = rs2  
# rd = success ? 0:1
```

- Further constraints (e.g., limiting number and type of instructions between the `lr/sc` instructions) guarantee that the sequence will eventually succeed.

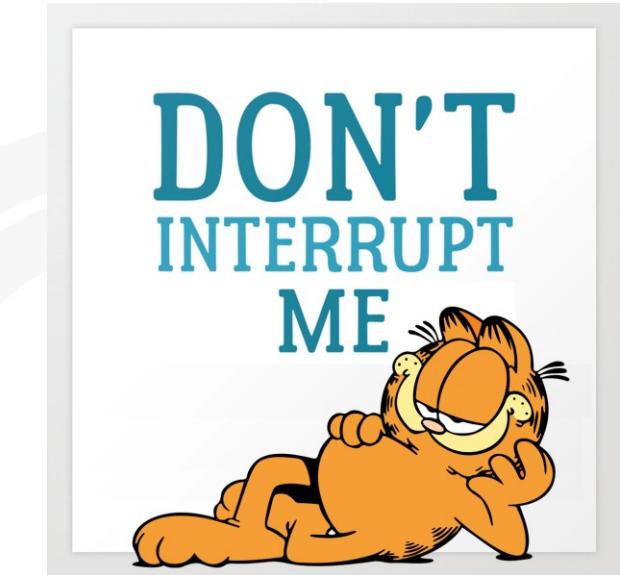
- RISC-V supports additional atomic instructions (e.g, `amoswap`, `amoadd`), but these can be implemented with `lr/sc` on simple processors.

Atomic Swap implemented with LL/LR

```
again:                                // swap x23 with (x20)  
    lr.w x10, (x20)      // x10 = *lock  
    sc.w x11, (x20), x23 // x11 = succeed/fail?  
    bne x11, x0, again   // spin if store failed  
    addi x23, x10, 0      // x23 = loaded value
```

Another approach: Disabling interrupts

- **Disabling interrupts** is a simple solution for mutual exclusion by ensuring there are **no context switches** while accessing the critical section.
- Problems with disabling interrupts:
 - Only available to the kernel
 - Can't allow user-level to disable interrupts!
 - Insufficient on a multiprocessor
 - Each processor has its own interrupt mechanism
 - “Long” periods with interrupts disabled can wreak havoc with devices
 - Use disabling of interrupts to build higher-level synchronization constructs



```
struct lock { }
void acquire(lock) {
    cli(); // disable interrupts
}
void release(lock) {
    sti(); // reenable interrupts
}
```

Final point... Semaphores!

- I know, since the beginning of this lecture, you've been waiting for me to talk about **semaphores**...

- Teman's First Law of Semaphores:**

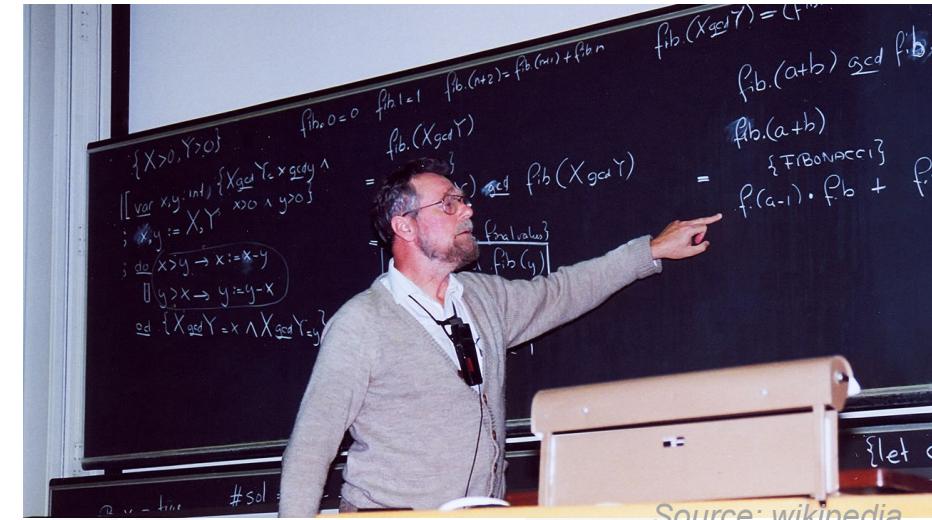
- The only thing anyone ever remembers from their undergrad operating systems class is that they learned about semaphores.

- Teman's Second Law of Semaphores:**

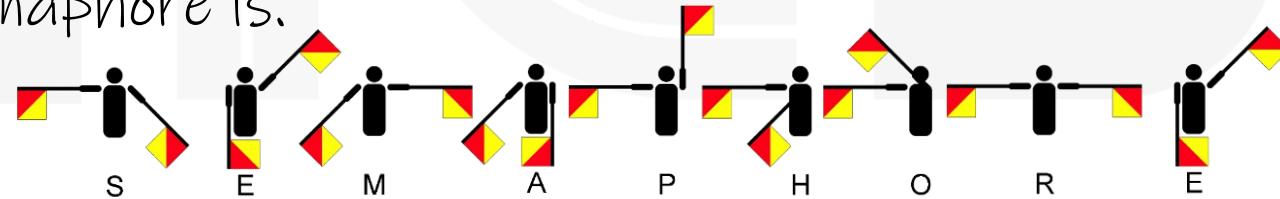
- No one actually remembers what a semaphore is.

- Teman's Third Law of Semaphores:**

- The only thing anyone remembers about semaphores is that Dijkstra invented them and there were some variables that come from Dutch.



Source: wikipedia



So, what is this Semaphore thing?

- Basically, it's a **shared, non-negative integer**.
- You mean, it's a **variable** that can have values, such as 0, 1, 2, etc.?
 - Yes... that's really all it is.
- So, what's the big fuss? Why is it different from any other variable?
 - Well, it has two operations that can be applied to it: **wait** and **signal**.
 - **wait(sem) (P)**: Block until **sem>0**, then subtract 1 from **sem** and proceed
 - **signal(sem) (V)**: Increment **sem**
 - Both of these operations must be atomic!
- And how exactly does that help us?
 - It allows one thread to **signal** to another.
 - Signaling is **non-blocking** (unlike acquiring a lock), which is very useful.

P is for “Proberen”
("test" or "try")
V is for “Verhagen”
("increase")



Signaling with Semaphores

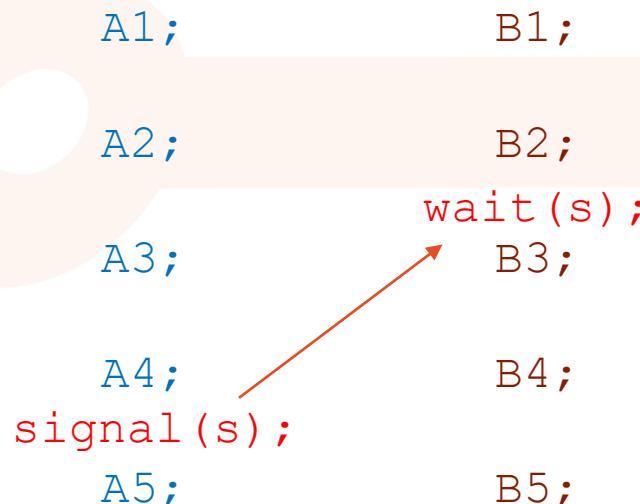
- **Example:** Semaphores for Precedence

- Two processes: A and B
- Goal: statement A4 should complete before B3 begins

```
semaphore s=0
```

Process A

Process B



- **Example:**

Bounded Buffer with Semaphores

SHARED MEMORY:

```
char buf[N]; // The buffer  
int in=0, out=0;  
semaphore chars=0;
```

PRODUCER:

```
send(char c) {  
    buf[in]=c;  
    in=(in+1)%N;  
    signal(chars);  
}
```

CONSUMER:

```
char rcv() {  
    char c;  
    wait(chars);  
    c=buf[out];  
    out=(out+1)%N;  
    return c  
}
```

* Insufficient for multiple producers and consumers. Requires mutex!

So, let's get something straight

- **Semaphores are not Mutexes!**

- A **Mutex** (a lock) is used for **mutual exclusion** to protect a **critical section**.
- A **Semaphore** is (remember those flags?) used to **signal a process**.

- So why is everyone (*I mean EVERYONE*) so confused?

- Because someone, somewhere decided to confuse everyone by calling a Mutex a “**binary semaphore**”.
- Because you can implement a Mutex with a Semaphore.
- And because certain systems (e.g., FreeRTOS) only provide semaphores.
- And just to help you remember, lets see what happens if we use semaphores at our tactically situated gas station...

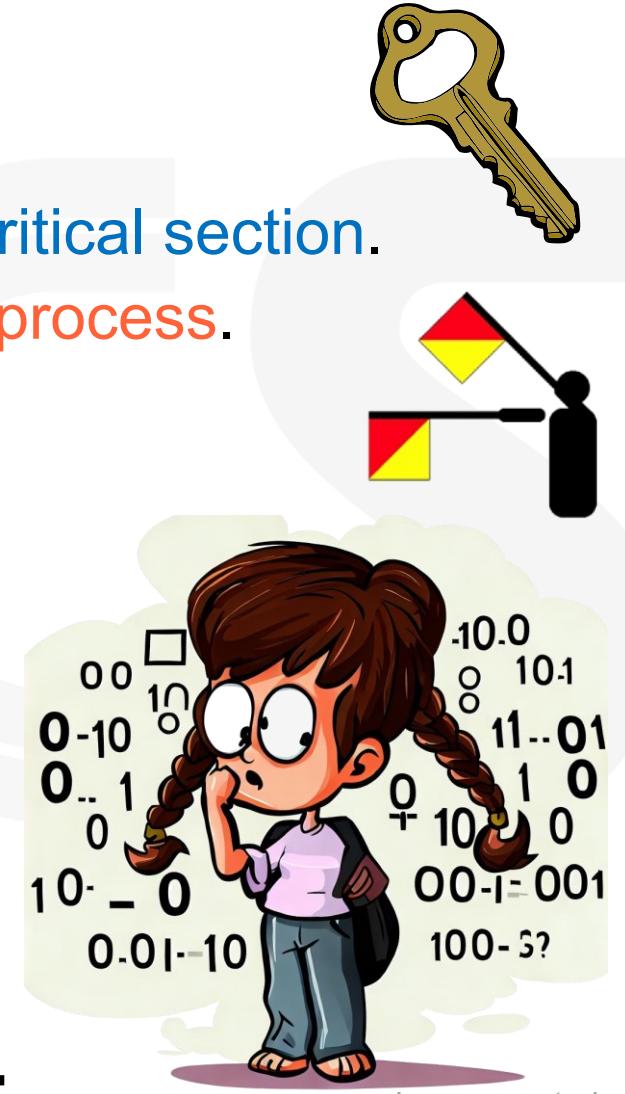


Image created with
Bing Image Creator

Gas Station Analogy: Semaphore

- Gas station has several bathrooms and several keys.

- First customer acquires key.
- Second customer acquires key too.
- Third customer waits for others to finish.
- First guy finishes and releases the key.

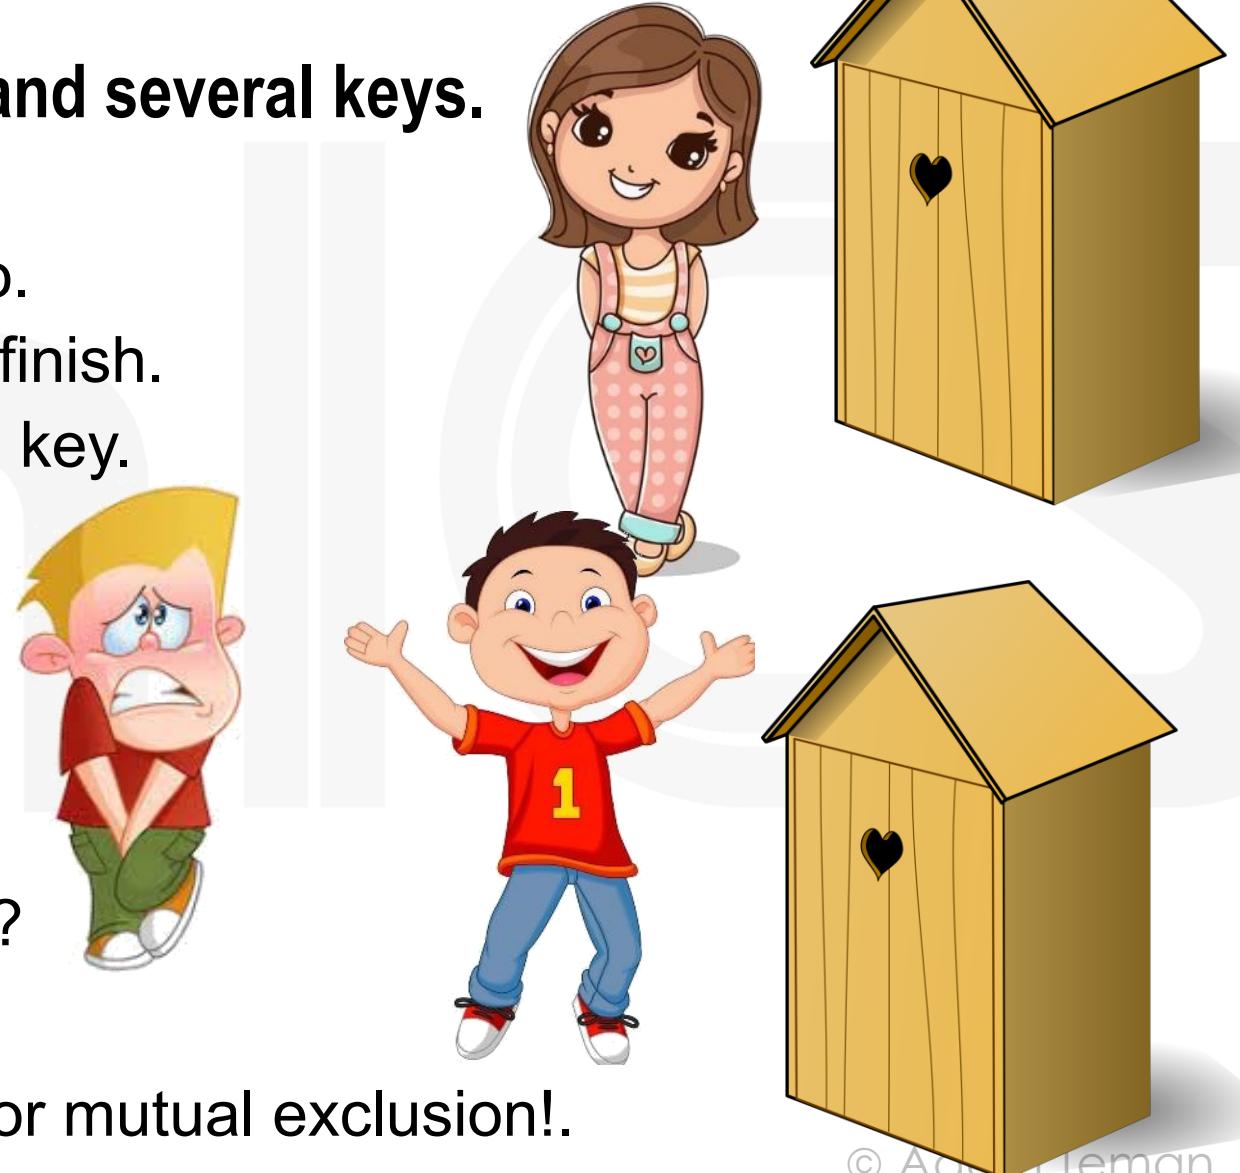


- **But...**

- How does the waiting customer know which bathroom is occupied?

- **Remember**

- Semaphores should not be used for mutual exclusion!.



Introduction

The Kernel

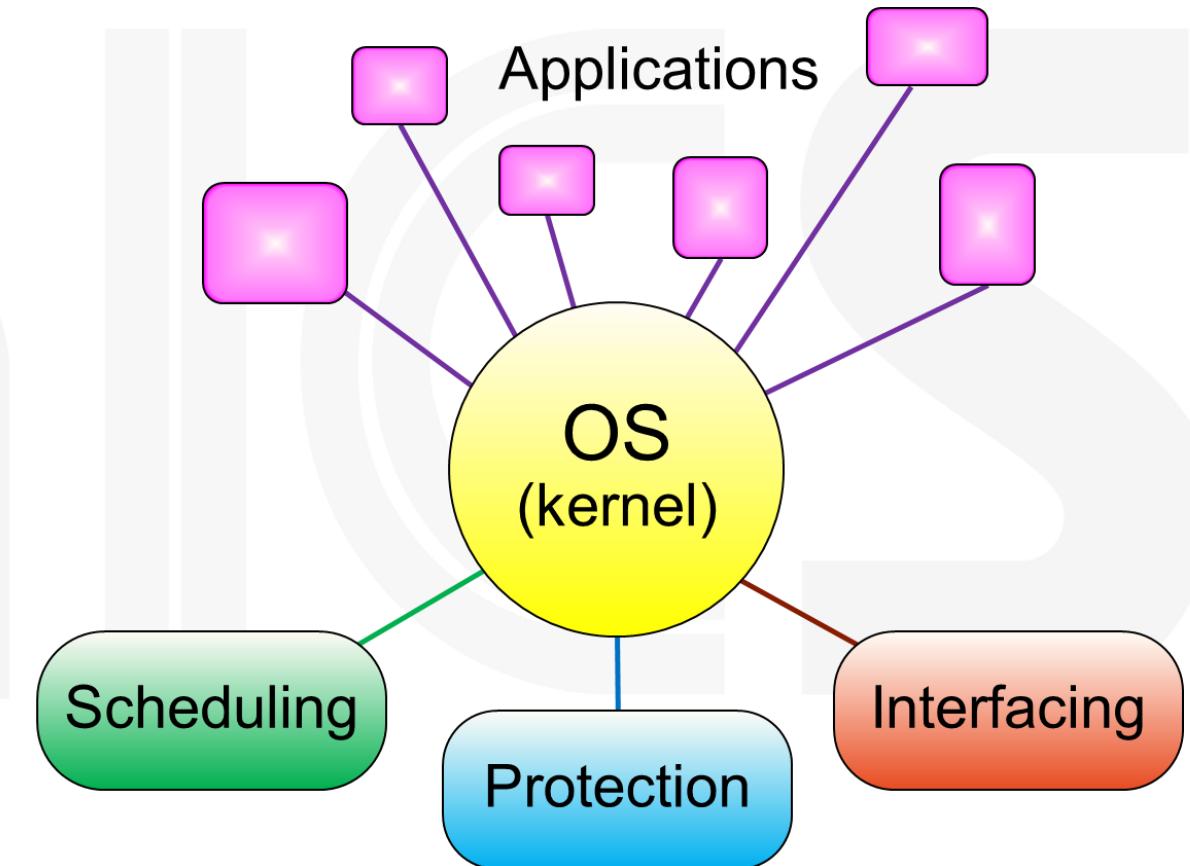
Interrupts

Scheduling

Synchronization

RTOS

RTOS



The Alexander Kofkin
Faculty of Engineering
Bar-Ilan University



Do we need OS for Embedded Systems?

Not always!

- For example, we want to **multitask several programs**

- The *cyclic executive* approach:

- **Pros**

- Simple implementation
- Low overhead
- Very predictable

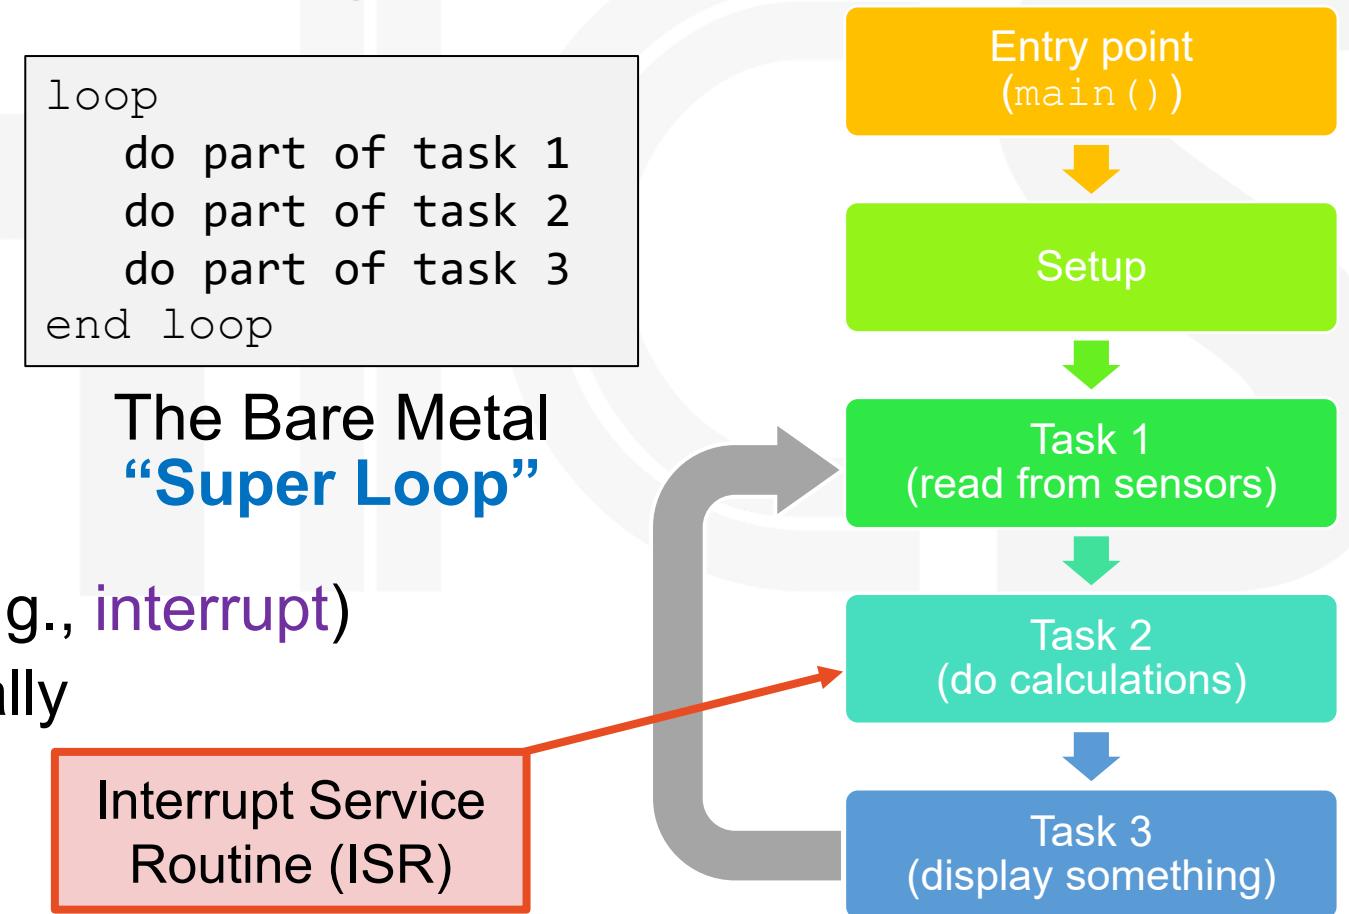
- **Cons**

- Can't handle sporadic events (e.g., *interrupt*)
- Code must be scheduled manually

```
loop  
    do part of task 1  
    do part of task 2  
    do part of task 3  
end loop
```

**The Bare Metal
“Super Loop”**

Interrupt Service
Routine (ISR)



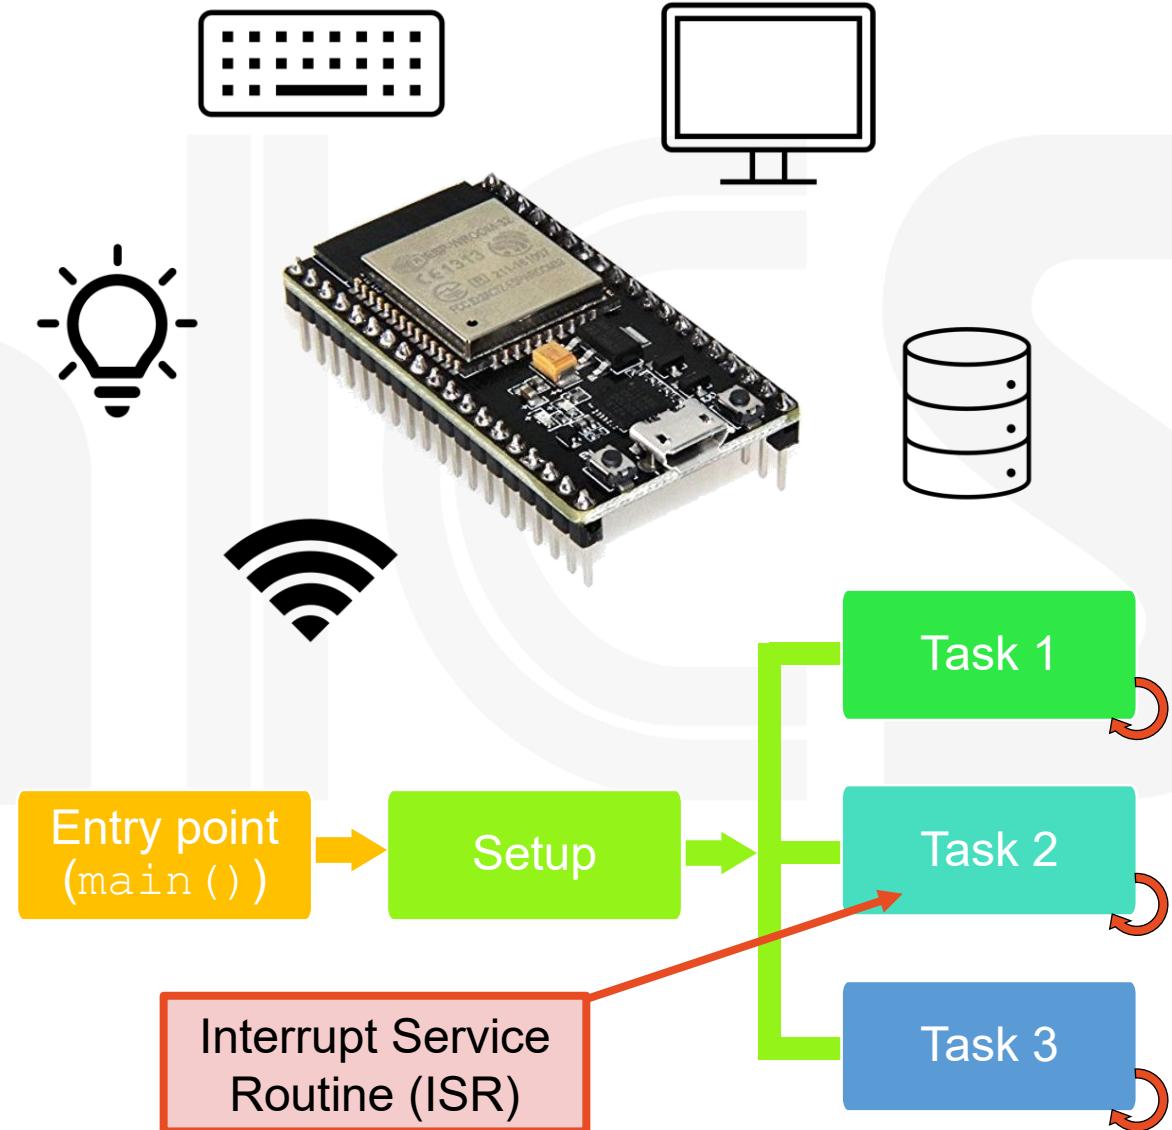
So, when do we need an OS?

- If we have a more complex system...

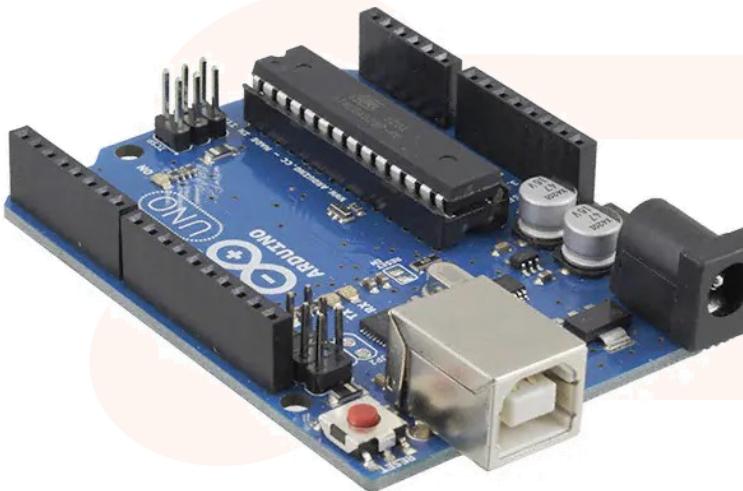
- Embedded systems need an OS for:

- Convenience
- Multitasking and threads
- Real-time requirements

- This is often provided by a **Real Time Operating System (RTOS)**

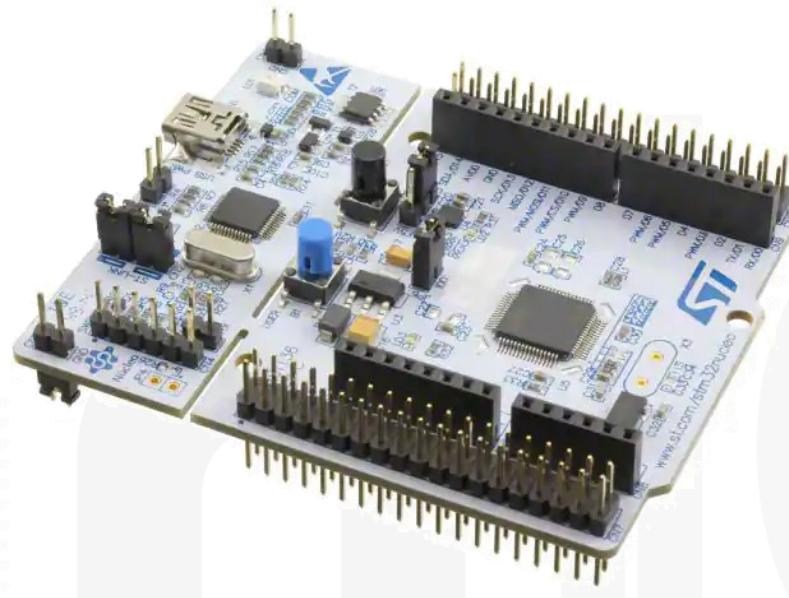


The choice is (also) hardware dependent



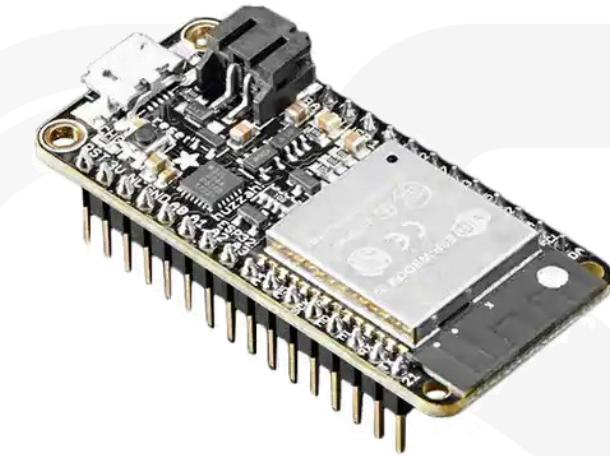
ATmega 328p

- 16 MHz
- 32 kB flash
- 2 kB RAM



STM32L476RG

- 80 MHz
- 1 MB flash
- 128 kB RAM



ESP-WROOM-32

- 240 MHz (dual core)
- 4 MB flash
- 520 kB RAM

Super Loop



RTOS

Source: Shawn Hymel, Digikey

© Adam Teman, 2023

Why “Real Time”?

- A **general-purpose operating system (GPOS) facilitates:**
 - Management of resources (CPU, memory, peripherals)
 - Multi-tasking (scheduling) and protection (synchronization)
 - User interface
- A GPOS can (usually) schedule processes based on **priority**
 - This enables differentiation between **more critical** and **less critical** tasks
- **However, timing cannot be guaranteed with standard scheduling policies**
 - A GPOS may switch from a critical task to another operation, completing the critical task only when the opportunity arises.
- **“Real Time” requirements mean that a task must be completed within a specific timeframe (usually within milliseconds).**



Image created with
Bing Image Creator

Implications of Real Time operation

- In a real-time system, a process has the following requirements:

- **Initiation Time:**

The time at the process goes from the **waiting** to the **ready** state

- **Deadline:**

The time at which the computation **must be finished**.

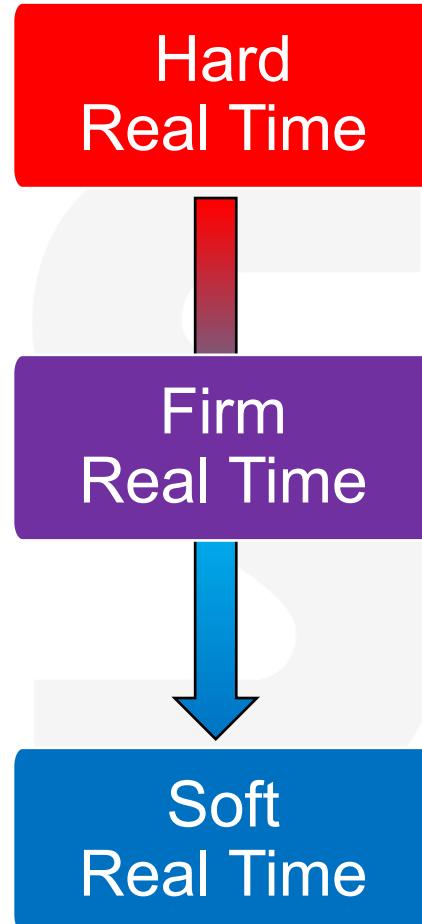
- What happens when a process misses a deadline?

- **Catastrophic**, e.g., missed deadline in automotive control system

- **Impede Quality**, e.g., missed deadline in telephone system

- For **safety-critical systems**, compensatory measures will be taken, e.g., approximating data or switching to a special safety mode.

- For **non-safety-critical systems**, the failure may be ignored or maybe bad data will not be propagated (i.e., silence in an audio stream)



So how do we achieve “Real Time”?

- Easy...

Priority-based Preemptive Scheduling!

- Split and prioritize the application code

- Each task is designed separately
- Each task is given a priority
- Tasks generally wait for an event to trigger
- The RTOS always runs the highest priority task that is available

- RTOSs are preemptive

- Most RTOS have a time-based interrupt (“tick”)
- Every tick, the running task is interrupted (preempted), the scheduler is invoked and it decides which task should run next.

Easy to program

Prioritize the most important task

Event-triggered – avoids polling

Low-priority tasks don't affect the responsiveness of high priority tasks

Preemptive Scheduling

ISR
(H/W)

Priority
(S/W)

1

0

OS

Task A

Task A

Task B

Task C

Task C

Task A

Idle

“tick”
1 ms

time

```
void EachTask (void){  
    Task initialization;  
    while (1) {  
        Setup to wait for event;  
        Wait for MY event to occur;  
        Perform task operation;  
    }  
}
```

```
void ISR (void){  
    Entering ISR;  
    Perform Work  
    Signal or send message to task;  
    Leave ISR;  
}
```

GPOS vs. RTOS

- GPOS

- Complex for general-purpose – multiple applications on a system
- Maximize resource usage, optimize average run time, ensure fairness
- Non-predictable behavior
- Pre-emptive or non-Pre-emptive
- Coarse timer resolution (~10ms)
- Large, protected virtual memory
- Heavy and Large



- RTOS

- Simple for embedded computing – single application on a system
- Ensure timing, meet hard deadlines
- Predictable
- Pre-emptive and priority-based
- Fine timer resolution (~1ms)
- Small, flat memory hierarchy
- Lightweight and small



Drawbacks of Using an RTOS

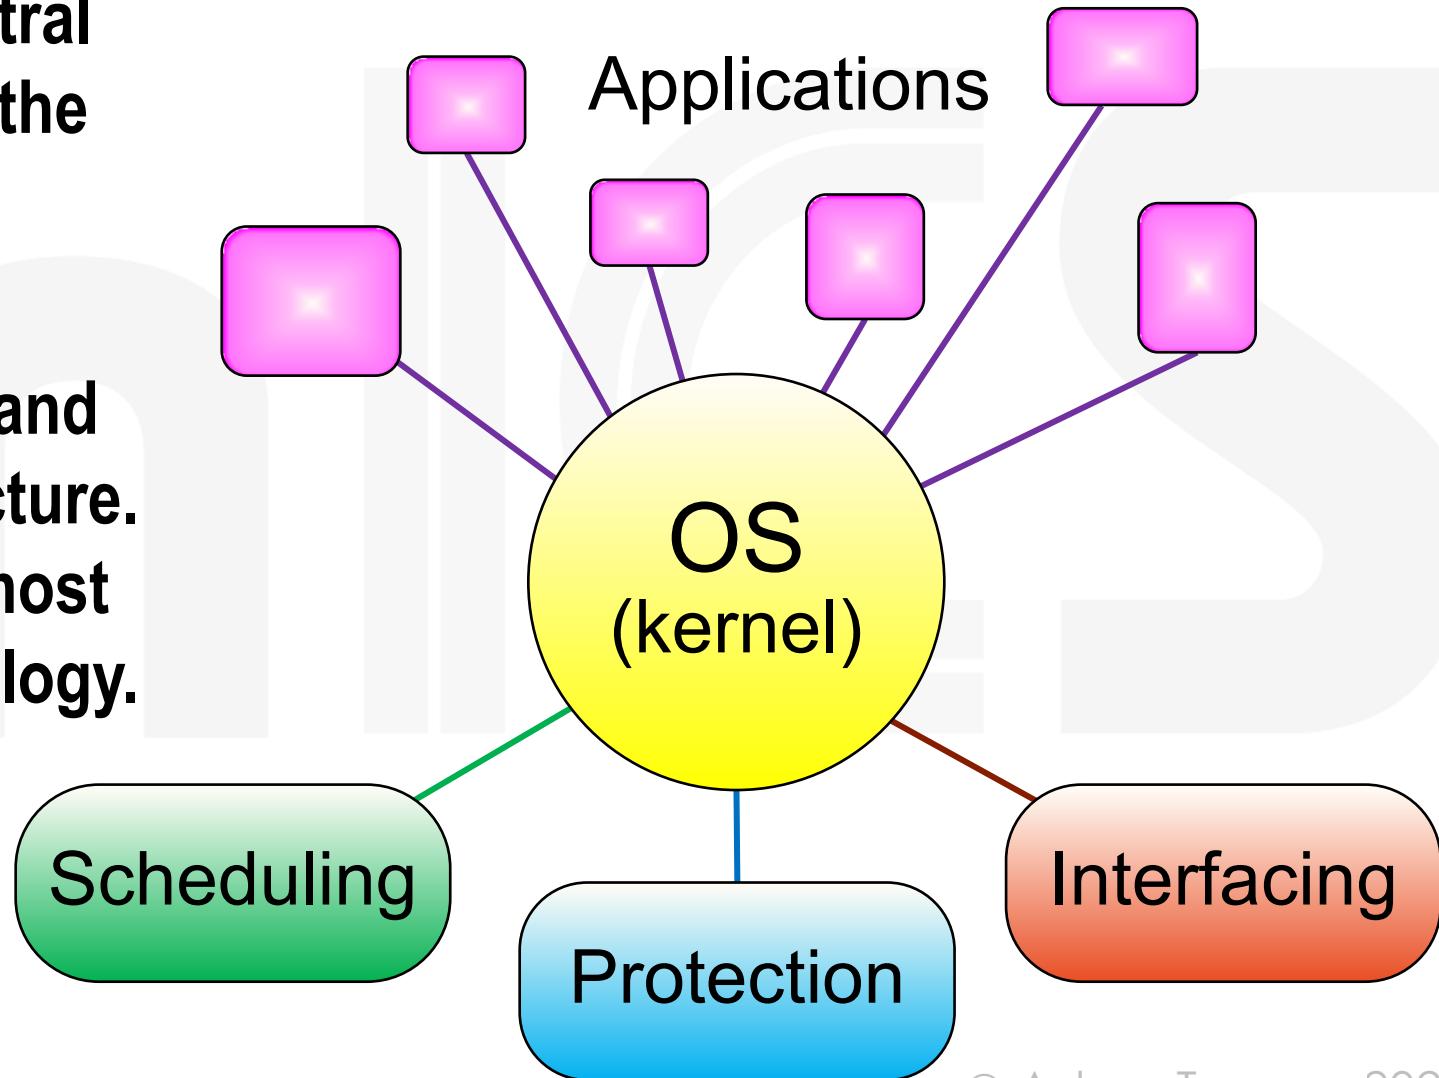
- The RTOS itself is code and thus requires **more Flash**
 - Typically, between 6-20K bytes
- An RTOS requires **extra RAM**
 - Each task requires its own **stack**
 - The size of each task depends on the application
 - Each task needs to be assigned a **Task Control Block (TCB)**
 - About 32 to 128 bytes of RAM
 - About 256 bytes for the RTOS variables
- You have to **assign task priorities**
 - Deciding on what priority to give tasks is not always trivial
- The services provided by the RTOS **consume CPU time**
 - Overhead is typically 2-10% of the CPU cycles
- There is a **learning curve** associated with the RTOS you select

Do You Need an RTOS?

- Do you have some **real-time requirements**?
- Do you have **independent tasks**?
 - User interface, control loops, communications, etc.
- Do you have tasks that could **starve other tasks**?
 - e.g. updating a graphics display, receiving an Ethernet frame, encryption, etc.
- Do you have **multiple programmers** working on **different parts** of your project?
- Is **portability** and **reuse** important?
- Does your product need additional **middleware components**?
 - TCP/IP stack, USB stack, GUI, File System, Bluetooth, etc.
- Do you have **enough RAM** to support **multiple tasks**?

Summary

- The operating system is the central piece of software that connects the programmer to the hardware.
- Operating system practicalities and research go well beyond this lecture. But we've learned some of the most important concepts and terminology.



Main References

- Gribble, Lazowska, Levy, Zahorjan, CSE 451, University of Washington
- Patterson, Hennessy “Computer Organization and Design – The RISC-V Edition”
- Berkeley CS-61C
- Null, Lobur, “The Essentials of Computer Organization and Architecture”
- Silberschatz, Gagne, Galvin, "Operating System Concepts, Ninth Edition"
- Venkatasubramanian, UC Irvine ICS 143
- Waterman, Asanovic, “The RISC-V Instruction Set Manual”
- Marilyn Wolf, “Computers as Components”
- UMASS CS-377
- Silicon Labs “Everything You Need to Know about RTOSs in 30 Minutes”
- Shawn Hymel, DigiKey “Introduction to RTOS”